

NAVIGATIONAL CYBERNETICS 2.5 MxBv, POSITIONING SERIES.

THE BODY THEY ARE BUILDING — PART III

THE URGRUND LABORATHEORY

WHITEPAPER V.00.01, 2026

The Body They Are Building — Part III. The Urgrund Laborattheory Whitepaper v 00.01

Navigational Cybernetics 2.5 — White Paper on the Research Programme

What the corpus investigates, which objects it distinguishes, and what architecture it builds from them

Maksim Barziankou (MxBv)

Navigational Cybernetics 2.5 · The Urgrund Laborattheory

Poznań · 2026

Status: working synthesis of the research programme.

Corpus slice: 22 July 2026.

Basis: 150 fully read and mapped works, the current version of the corpus map, and the active formal branch *Non-Causal Non-Causality*.

DOI: [10.17605/OSF.IO/WXPHF](https://doi.org/10.17605/OSF.IO/WXPHF).

License: CC BY-NC-ND 4.0.

Abstract

Navigational Cybernetics 2.5, or NC2.5, is a research programme about long-lived adaptive systems for which acting successfully at each separate moment is not enough. Such systems must preserve the possibility of continuation, their own identity, the provenance of their authority, and internal coherence along a trajectory that passes through irreversible changes.

The programme's initial question is simple:

What must an architecture contain so that a system can change, pay for change, move between regimes and layers, yet not substitute a sequence of locally successful actions for the continuation of its existence?

The corpus builds its answer not around a single quantity. It separates at least eight functions: history and carrier, observation, admissibility geometry, selection and authority, load and internal time, meaning and liveness, coordination and evidence, and the constitution of a system out of components. The central verdicts are likewise separated: does the system continue to act; does it remain viable; is the identity

witness preserved; does the transition remain admissible; does the acting object hold the prior authority. None of these answers follows automatically from the others.

The corpus's mathematical language combines finite history spaces, factorisation and observation fibres, information theory, viability and reachability, Hodge / de Rham / cohomological constructions, categorical forgetful separations, typed authorities, and deterministic executable checks. These tools are for the most part known. The programme's principal originality lies in their composition and in the strict separation of operative roles. The primitive-level claim is kept narrow and object-specific: it concerns the non-causal trajectory predicate and the structural-pair spin layer $([\eta_S], d\omega_S)$, explicitly separated from the period / magnitude ledger readouts. Double Fibre is claimed at the level of the architectural synthesis of two typed forgetting operations, not as an invention of factorisation or conditional entropy.

The corpus does not claim that all adaptation must follow one dynamics, that any change is degradation, that any constraint is new, that spin is identical to identity, or that future-indexed geometry means reverse causation. All strong results are indexed by a declared class of systems, carrier, observation channel, horizon, predicate, and authority.

1. Subject of the investigation

Most contemporary systems are evaluated by local units:

- correctness of an answer;
- quality of a decision;
- reward;
- error;
- compliance with a rule in the current state;
- completion of a single transaction;
- match between observed behaviour and what was expected.

For a short horizon this is often enough. For a system that must act for months or years, change its own rules, transfer state between carriers, interact with other agents, and preserve the right to continue its prior line, local evaluation ceases to be complete.

NC2.5 investigates a different object: **the existence of a system on a trajectory**.

In this formulation a state is only a cross-section. It may show that the system is working now, but it does not answer:

- how it arrived at this state;
- which admissibility criterion held along the way;
- whether the identity-bearing witness is preserved;
- how much irreversible load has accumulated;
- which future continuations are still available;
- whether the present executor holds the prior right to act;
- whether the observed system is a continuation of the prior object or only a functionally similar reconstruction of it.

The corpus's principal thesis therefore has the form:

Identity is a trajectory predicate inside drift.

Identity is neither a value at a point nor a sum of deeds. It is a predicate over the history of transport. Drift, in turn, does not reduce to error: it is the medium in which irreversible changes take place, and navigation is motion through this medium that preserves the declared witness, admissibility, and possibility of continuation.

2. What this white paper does and does not do

The document performs four tasks.

1. It gathers 150 mapped works and the latest formal branch into one research programme.
2. It shows which objects are load-bearing and which serve as bridges, applications, or phenomenological registers.
3. It separates established finite results from theorem programmes, engineering declarations, and open hypotheses.
4. It connects the research corpus to the PETRONUS product line without turning the products into proofs of the theory.

This text does not replace:

- the axiomatic core NC2.5 v2.1;
- the formal companion papers;
- the deterministic harnesses;
- the corpus map;
- the source articles;
- external review and the publication seal.

Nor does it raise the evidential status of the early works. If a source text is an essay, a programmatic declaration, an application note, or an internal record, here it remains exactly that. Strong ideas from the early corpus are used only after typing and only where an explicit bridge to the formal layer exists.

3. How the research method is built

3.1 Structure before name

The programme develops from form to mathematics. First a recurring structural distinction is discovered: action and the right to act; state and history; wear and identity; local admissibility and the fate of a regime; observation and enforcement. Then the object is decomposed into minimal non-mixable roles. Only after that is a mathematical language chosen.

The same meaning may therefore pass through several registers:

- phenomenological;
- philosophical;
- architectural;

- formal-mathematical;
- protocol;
- executable;
- product.

Coincidence of topic between registers is not a proof of their equivalence. A transition requires a bridge with declared types, carrier, channel, direction, and failure condition.

3.2 Typed separation

The corpus's main discipline is not to let one word perform several incompatible functions. In particular:

- observation is not enforcement;
- constraint is not selector;
- selector is not authority;
- history is not a memory snapshot;
- viability is not identity;
- liveness is not uptime;
- budget is not witness;
- amplitude is not winding;
- local vorticity is not global cohomology;
- coordination is not constitution;
- declaration is not causal production;
- present truth is not a cross-regime guarantee;
- reset is not restoration;
- public description is not implementation.

Most real failures in long-lived systems arise not from missing data but from the collapse of these types.

3.3 Class-relative claims

NC2.5 does not claim a universal physical theory of all systems. A result holds for a declared class once the following are given:

- the carrier and the system boundary;
- the space of histories;
- the class of dynamics and disturbances;
- the horizon;
- the admissibility criterion;
- the identity target and its witness;
- the model of load and recovery;
- the observation channel;
- the authorities of proposal, selection, enforcement, and revision.

Without these data a phrase remains an architectural intuition, not a theorem.

3.4 Countermodels before universality

The corpus uses short countermodels as its principal way of holding the boundaries of its claims. If two complete cases give the same visible trace but different target verdicts, the visible channel is insufficient for exact inference. If the same budget is compatible with different continuation geometries, budget does not determine continuation geometry. If the same amplitude is compatible with different topological witness states — in an explicit model with zero versus non-zero winding — an amplitude monitor does not certify the topological readout.

Such a countermodel does not prove that the target can never be known. It proves only the insufficiency of the particular declared channel.

3.5 Corpus as an auditable research object

The corpus stores not only results but also their provenance:

- exact snapshots of sources;
- SHA-256;
- edition diffs;
- the status of formulations;
- incompatibilities between early versions;
- bridge obligations;
- theorem-facing directions;
- open falsification surfaces.

The corpus map is therefore not a bibliography but the upper architectural layer. It does not erase failed early formulations but shows which later works refined, bounded, or refuted them.

4. Minimal architecture

Let h_t be the history of the system up to time t . For a long horizon a single state vector is insufficient. At minimum one must distinguish:

- $Live_t$ — whether the endogenous reinitiation of the declared live cycle is preserved;
- $Witness_t$ — whether the chosen identity witness is transported;
- Adm_t — whether the transition remains admissible relative to the declared predicate.

The corpus's summary typing schema records the preliminary verdict as

$$V_t = (Live_t, Witness_t, Adm_t).$$

In the corpus no coordinate automatically entails either of the other two. This is not a single theorem of the axiomatic core: the corresponding pairwise non-implications are established by different companion works and countermodels, and a concrete joint realisation requires a separate composition certificate.

- A system may remain operationally active yet lose the identity witness.
- The witness may be preserved while the budget is being spent.
- A transition may be admissible yet lead to a future loss of a non-trivial continuation interior.
- A live cycle may continue inside an object that is already unauthorised.

- A new object may be viable yet lack the prior authority.

Above this vector sit further objects:

- hidden generator hypotheses;
- phase state;
- remaining budget;
- available action set;
- observation and property channels;
- authority and revision records;
- constitution record;
- cross-layer carrier.

They do not become new coordinates merely because they have received a name. A standalone coordinate requires a separate target, an operational witness, and an independence argument.

5. The axiomatic core of NC2.5

5.1 Drift as medium, not error

In early control systems drift is usually understood as deviation from a reference. NC2.5 considers a broader class: irreversible deformation of identity-relevant structure in the course of interaction.

In this sense drift is not an error to be eliminated but the very medium of navigation.

This does not mean that every change is harmful. Drift can:

- preserve identity;
- change form while the witness is preserved;
- open new continuation regions;
- accumulate irreversible load;
- pass into rupture;
- lead to a separately typed repair or reconstitution; such a transition does not erase the prior burden history and is not a compensation of drift within the original seal.

Navigation therefore is not equal to the minimisation of drift. Its task is to lead the system through change while preserving the declared invariants and not passing off local smoothness as a proof of continuation.

5.2 Load and internal time

For the chosen axiomatic class the basic ledger entry has the form

$$\tau_{\text{rem}}(t) = C - \Phi(t), \quad \Phi(t) = \int_0^t \rho_{\Phi}(x(s)) \, ds, \quad \rho_{\Phi} \geq 0,$$

where C is the fixed structural capacity of the declared class, ρ_{Φ} is the instantaneous burden rate, $\Phi(t)$ is the monotone accumulated burden, and $\tau_{\text{rem}}(t)$ is the remaining viability budget. On the declared viability domain $\Phi(t) \leq C$ holds, hence $\tau_{\text{rem}}(t) \geq 0$; a state with $\Phi(t) > C$ is a terminal exit / a state outside the kernel, not an ordinary negative remaining budget.

From the single algebraic entry by itself the following do not follow:

- finite collapse;
- nesting of continuation sets;
- preservation of identity;
- monotone contraction of the whole continuation geometry.

These conclusions require separate premises. In the current core the constancy of $\mathbf{C}$ and the non-negativity of $\rho_\emptyset$ are precisely such separate class conditions, not consequences of subtraction. The corpus distinguishes active burden, passive burden, phase debt, and source persistence; repair, nonstationary capacity, or capacity growth belong to extensions or to a new sealed declaration and do not replace this core model. Structural pressure is especially important: load can grow even under zero task-level action.

Internal time is not a clock. It is a parameter of the remaining capacity to continue the declared class of motion. In various early works one symbol denoted both elapsed exposure and remaining capacity; the final architecture separates these types.

5.3 Admissibility before optimization

Optimization answers: which element to choose among those available. Admissibility answers: which continuations belong to the permitted class at all.

The architectural order has the form:

```
candidate generation
  → admissibility restriction
  → selection inside admitted support
  → enforcement
  → realised transition.
```

Hard admissibility is not a very large penalty. If a violation can be traded for sufficient reward, that is an objective trade-off, not a categorical exclusion.

At the same time, a Boolean gate by itself does not yet yield a protected architecture. Repeated membership queries, timing, rejection patterns, and explanation channels can make the boundary reconstructible. The corpus therefore types not only the verdict but also the query surface, the information access, and the revision authority.

5.4 Phase, regime, and false reversal

The same visible transition may have a different meaning in different phases of history. Phase Mechanics introduces the idea of unclosed phase debt: a locally correct contact may change future admissibility depending on whether it is a continuation, a commitment, a closure, or a rupture.

From this follows one of the corpus's stable distinctions:

Numerical return is not structural restoration.

A system may return to prior parameter values without restoring:

- the prior history;
- the prior witness;
- the prior authority relation;

- the prior continuation set;
- the prior seal.

A reset creates a new event. If it erases history and declares itself a return to the old object, that is a false reversal.

5.5 Non-actionability

The word non-causal in the corpus does not mean causal helplessness. A gate that blocks a transition causally changes the executed trace. The stable meaning is different:

- the ground of admissibility is not produced by the task objective;
- the ordinary optimizer has no right to rewrite the ground;
- the privileged predicate does not become a reward or a gradient;
- access to the boundary is limited by a declared channel contract;
- revision belongs to a separate authority.

A short formula:

Causal influence does not entail optimizer query, gradient, or write authority.

It is precisely the separation of access and authority, not the absence of a system-level effect, that is the load-bearing object.

5.6 Structural spin

On a fixed smooth Riemannian carrier (M, g) and in the Hodge-regular subclass, the dynamical entry has the form

$$F = -\nabla_g V + S,$$

where S is a chosen orthogonal non-gradient component. Before fixing the metric, boundary conditions, and Hodge representative, the formal entry has the gauge freedom $(V, S) \mapsto (V+u, S+\nabla_g u)$; outside the Hodge-regular subclass only the weaker conclusion "the dynamics is not a global gradient flow" is guaranteed.

S must not be automatically called identity, energy, residual cost, or winding. Its roles are defined by separate structural channels, readouts, and bridges.

Spin is needed by the programme as a language of preserved directedness, circulation, and witness-bearing structure. The exact separation claim is conditional: if a scalar energy/budget channel factorises through the gradient skeleton, it does not determine S ; in general any relation between energy and circulation requires a separately declared bridge and calibration. The same scalar account without such a bridge is compatible with different non-gradient geometry.

6. Identity: not a point, not a sum, not a budget

6.1 Trajectory predicate

Identity in the corpus is defined relative to a declared witness over a history. This may be:

- a topological class;
- a transportable covector;
- an anchor;
- a trace relation;
- a lineage certificate;
- a configuration / authorial continuity pair;
- another pre-typed witness.

Without a declared witness the question "is this the same object" is underdetermined.

6.2 Two ledgers

The work *Identity Does Not Drift* separates:

1. the **budget ledger** — how much the system has paid for switching, loads, and wear;
2. the **identity ledger** — whether the witness is preserved under transport.

A system may age, expend resource, and change channels yet preserve the witness. It may also remain viable and functionally continuous after the destruction of the witness.

The critical case is called:

survives-but-dead.

Here viability is preserved while the identity-bearing witness is extinguished. Operational continuity in such a case is not a proof of the continuation of the same authorised object.

6.3 Switching and tunnelled transport

Under channel switching, external similarity of states is insufficient. A transport condition on the witness-bearing locus is required. A parallel tunnel means that the traffic of switches passes through one part of the structure while the identity witness is carried along a protected locus that these operations do not alter.

This does not make switching free. The budget ledger pays. The claim is only that, under correct typing, the payment need not enter the identity coordinate.

6.4 Identity, admissibility, and authority

Three continuity claims must remain separate:

- **identity continuity** — whether the witness is preserved;
- **admissibility continuity** — whether the object remains admissible in the new regime;
- **authority continuity** — whether the prior authority transfers to the target.

Preservation of identity does not transfer authority automatically. Transfer of authority is a governance predicate that requires its own certificate. A change of carrier, transition map, or witness bridge opens

only change-dependent obligations; unchanged evidence may persist. But after the failure of an authorisation-bearing predicate, any further authority is a new certified grant, not a continuation of the old one.

7. The double fibre of verification

7.1 Two independent losses

Let $\mathcal{D}$ be a finite declared class of dynamic skeletons. For each $\delta \in \mathcal{D}$ there are finite sets of histories H_δ and a visible codomain Y_δ , an observation map $O_\delta : H_\delta \rightarrow Y_\delta$, a finite declared family of admissibility selectors $\mathcal{A}(\delta) \subseteq \{0,1\}^{H_\delta}$, and a compatibility relation $C_\delta \subseteq H_\delta \times \mathcal{A}(\delta)$. The corpus singles out two forgetting operations.

1. Indexed observation forgets lineage:

$$O_\delta : H_\delta \rightarrow Y_\delta.$$

2. The forgetful projection U forgets the admissibility selector:

$$U : (\delta, \alpha) \mapsto \delta, \quad \alpha \in \mathcal{A}(\delta), \quad \alpha : H_\delta \rightarrow \{0,1\}.$$

The full declared space and the visible map have the types

$$\begin{aligned} \Theta &= \bigsqcup_{\delta \in \mathcal{D}} \{\delta\} \times C_\delta, \\ Z(\delta, h, \alpha) &= (\delta, O_\delta(h)), \\ F_2(\delta, y) &= Z^{-1}(\delta, y). \end{aligned}$$

The index δ does not permit silently gluing identically looking artifacts of different skeletons. The same visible value within one typed fibre may have different compatible histories and selectors; their joint inverse image forms the **double fibre of verification**.

7.2 Factorisation criterion

For a finite declared class, a target Q is exactly computable from the visible representation Z if and only if Q is constant on every double fibre. Under a full-support prior this is equivalent to:

$$H(Q \mid Z) = 0.$$

Positive uncertainty of the hidden pair need not mean ambiguity of the target. Several hidden histories may still give one verdict. The corpus therefore separates:

- hidden-pair binding;
- target ambiguity;
- reconstruction;
- property detection.

7.3 Three levels of knowledge

The epistemic ladder has three rungs.

1. **Exact verdict:** the target is constant on every fibre.
2. **Statistical property detection:** the conditional distributions differ and a calibrated positive property channel exists.
3. **Hidden-state reconstruction:** the provenance or the full hidden state is recovered.

The second rung is possible when the third is impossible. This is the central result of the *Identifiability Bridge*: an NR-bound does not mean that no property can be detected; a positive channel requires its own Bridge Detectability Premise, an ensemble, and calibration.

7.4 Constructive repair

Refinement is compared on partitions of the space $\Theta: \mathcal{O}' \geq \mathcal{Z}$ if equality of $\mathcal{O}'$ -values always entails equality of $\mathcal{Z}$ -values. For a finite target the least sufficient visible refinement has the form

$$\mathcal{Z}_Q = (\mathcal{Z}, Q).$$

It refines $\mathcal{Z}$, makes Q computable, and has a universal property: if $\mathcal{O}' \geq \mathcal{Z}$ and $Q = f \circ \mathcal{O}'$, then $\mathcal{O}' \geq \mathcal{Z}_Q$. Under a declared full-support prior its exact information price equals $H(\mathcal{Z}_Q|\mathcal{Z}) = H(Q|\mathcal{Z})$, and any sufficient $\mathcal{O}'$ must add at least the residual target entropy.

This does not mean the target must be literally disclosed to the optimizer. The construction fixes the minimal information content of a sufficient channel; the physical instrument, the right to read, and the right to use remain separate architectural obligations.

7.5 Verification is not enforcement

An observer may classify correctly and block nothing. An enforcer may block without knowing the full hidden state. For a finite typed event alphabet a deterministic reference monitor is sufficient only under the full set of premises: the declared decidable prefix-closed safety property S contains the empty prefix; the monitor holds an authoritative committed prefix p ; on each proposed event e the exact S -monitor commits e if and only if $pe \in S$, otherwise it rejects the event without changing p ; every committed event is visible to the monitor; every commit passes only through the monitor, with no bypass; the monitor evaluates the authoritative prefix; the membership procedure is correct. Then the monitor prefix coincides with the authoritative system prefix and remains in S . But liveness, repair, and recovery require additional contracts.

Thus verification, reconstruction, property detection, enforcement, revision, and coordination are not one score.

8. Structural spin: one carrier, two channels, three meters

8.1 Operational spin

The early work *Operational Spin* introduces a local rotational residue as the antisymmetric part of a gradient / Jacobian construction. This is a local differential object: it detects non-gradient structure in a neighbourhood.

8.2 Structural spin

On a fixed (M, g) set $\omega_S = S^\perp \{b_g\}$. Since an arbitrary ω_S need not be closed, the cohomological channel requires a declared selector

$$P_{cl} : \Omega^1(M) \rightarrow Z^1(M), \quad \eta_S = P_{cl}(\omega_S),$$

which is a fixed linear Hodge projection onto the closed summand $\text{im } d \oplus \mathcal{H}^1_g$ for declared metric, boundary conditions, and regularity class: it is idempotent, annihilates the co-exact summand, and fixes every closed input. Hence $P_{cl}(\omega_S + du) = P_{cl}(\omega_S) + du$, in particular $[P_{cl}(\omega_S + du)] = [P_{cl}(\omega_S)]$; at the same time $d(\omega_S + du) = d\omega_S$. The selector is part of the type; the minimal carrier claim concerns (M, g, F, P_{cl}) , not the bare formula $F = -\nabla V + S$.

The companion structural signature separates two inequivalent structural channels:

$$C_{str}(S; P_{cl}) = ([\eta_S], d\omega_S).$$

The $d\omega_S$ half yields the operational spin. The global class is read through periods and winding witnesses. This structural pair dissociates in both directions:

- $d\omega_S$ may be non-zero while $[\eta_S]$ is zero;
- $[\eta_S]$ may be non-zero while $d\omega_S$ is zero.

It must not be conflated with the Layer-3 ledger pair $C_{led} = (\sigma_H, \Lambda_{loc})$, where σ_H is a period readout and Λ_{loc} is a magnitude / norm readout on one declared locus. For the ledger pair the triangle inequality gives only $\sigma_H \neq 0 \Rightarrow \Lambda_{loc} > 0$; the converse is false. The two-sided dissociation of the structural pair does not license two-sided independence of the ledger readouts.

8.3 Three readouts

The current typing refinement distinguishes three scalar readout families:

$$\begin{aligned} W_\gamma(\eta_S) &= \oint_\gamma \eta_S, \\ A_\gamma(\alpha) &= \int |\alpha(\gamma(t))| dt, \\ Q_P(\omega_S) &= \int_P d\omega_S. \end{aligned}$$

Here $\gamma \in Z_1(M)$ is a smooth closed 1-cycle, $\eta_S \in Z^1(M)$ and $\alpha \in \Omega^1(M)$. Hence W_γ is a de Rham period / winding readout; the integral of an arbitrary non-closed α over a loop is a circulation but not a cohomological period. A_γ is a path total variation (a seminorm), not an arbitrary amplitude, and Q_P is a vorticity-flux readout. For a globally defined ω_S on a closed 2-cycle P the last integral is zero by Stokes; a non-zero flux requires $\partial P \neq \emptyset$ or a separately declared singular / patchwise construction.

Apart from these three readouts stands a fourth object — the dynamic residual used in a drift or maintenance ledger.

Mixing structural channels, scalar readouts, and the residual makes an empirical result unfalsifiable. Every bridge must declare a probe and an observation map.

8.4 Identity bridge

A topological witness is not identity automatically. It becomes an identity witness only through a separate bridge declaration. A local vortical channel detects non-gradient structure but by itself likewise does not

certify a transfer of identity.

Therefore a collapse of winding under preserved amplitude means only a readout dissociation and the incompleteness of an amplitude-only monitor; structural-channel dissociation additionally requires a separate measurement of $d\omega_S$. The survives-but-dead verdict also requires:

- an independent viability / continuation control;
- a pre-declared identity-witness bridge.

8.5 Empirical surface

The work on cortex-wide rotating waves uses a 2π phase winding around a centre. This is not the same H_1 generator, because the physical cortical space and the abstract state-space carrier differ. What overlaps is the type of invariant: a degree-one winding around a singularity.

The corpus's falsifiable stake is narrow: at amplitude-matched epochs a preregistered loop-local winding index must be able to collapse separately from the 2–8 Hz amplitude. One loop-local verdict does not prove the collapse of the whole class without a generator-complete loop family. If, after correct controls, such a dissociation is impossible, the empirical reach of the interpretation contracts.

9. Future-indexed admissibility environments

9.1 The architectural inversion

An ordinary planner acts in the order:

```
present → proposed action → predicted future → evaluation.
```

Non-Causal Non-Causality investigates the reverse order of architectural dependence:

```
history now
  → future-indexed admissible environment
  → present projection
  → present action.
```

The first decision object becomes not the action but the set of reachable continuations satisfying trajectory identity, admissibility, and horizon-relative viability. An action enters the existential support $\Pi^{\exists}$ if and only if it is the first coordinate of at least one such continuation; this is the weakest assurance mode and does not mean robustness or liveness. Actual runtime admission uses a pre-declared mode-specific support Π_t^{act} — existential, adaptive-robust, or recursively seal-relative robust — and authenticated enforcement. The task objective may rank only an already-formed active support, but obtains no authority to create or rewrite its ground.

9.2 This is not retrocausality

The future-indexed set is built in the present from information available in the present. No physical event from the future enters the input. The future index describes a position inside a formal continuation, not a direction of information flow.

Non-causality here means a relation of grounding and authority:

- the future environment is not a reward;
- the privileged generator is separated from the task optimizer;
- the consumer receives a typed readout;
- enforcement is causally effective;
- revision opens a new versioned certificate.

9.3 History-generated recurrence

The temporal joint has the form:

```
(history_t, specification_t, seal/provenance_t)
  → independently generated future environment_t
  → mode-specific active support_t
  → proposed action_t
  → authenticated one-shot decision_t
    bound to (proposal, history_t, assurance context_t,
              specification_t, time, horizon, regime, version)
  → mediated realised event_t
  → exact append to history_{t+1}
  → next authenticated observation uniquely interpreted as history_{t+1}
  → valid successor specification/provenance_{t+1}
    binding the prior seal specifically to realised event_t
  → independently recomputed future environment_{t+1}.
```

History alone does not generate the future environment: an active specification, its provenance, a validity envelope, and confined generator inputs are required. Every arrow of the recurrence is a typed binding and fails closed if it is undefined or does not pass authentication: the decision is bound to the whole declared context, the successor observation is uniquely interpreted as a new history prefix, and the successor provenance binds the prior seal specifically to the realised event, excluding replay and substitution. The present, having become a mediated event, enters the past; the extended history participates in the independent construction of the next future environment. This is not a temporal circle but a chain of strictly growing prefixes. The prior seal is not rewritten: a change of specification creates a new certificate bound by lineage to the previous one and cannot retroactively rescue an old verdict.

9.4 Atmosphere

In the platform language the object is called the atmosphere of admissibility. It is neither a single moral table nor a prediction. Each person or system supplies constitutive content, and the platform preserves:

- provenance;
- authority separation;
- query confinement;
- cross-temporal coherence;
- version lineage.

The platform does not hand everyone one future. It provides machinery for building and living one's own accountable space of admissible futures.

10. Operator-relative chaos and Point of Attention

This branch remains partly unpublished, so only the large-grain meaning is fixed here.

10.1 Chaos and order as verdicts

The Theory of Internal Chaos treats chaos and order not as opposite absolute properties of a substrate but as two operator-relative verdicts over one geometry.

An operator has a bounded horizon:

- what it is able to hold as a whole;
- within what it is able to orient;
- which load it is able to distinguish;
- which continuation relations are available to its observation and action architecture.

Personal order is not the absence of motion. It is the remaining margin up to the horizon of holding and orientation. Therefore:

- more bandwidth does not mean a farther architectural horizon;
- noise and motion are not identical to chaos;
- stillness does not certify order;
- a complex system may be mobile and internally held;
- a static system may be operator-incoherent.

10.2 Attention as readout and intervention

The corpus separates two models.

- **Flashlight:** attention changes only what is read.
- **Sculptor:** attention enters the update law and changes the future representation state.

They can be distinguished by an intervention protocol: under the same prior history and input, matched interventions on attention must yield a difference in a later standardised probe that cannot be routed through a readout-only path.

10.3 Point of Attention

Point of Attention is a candidate instrument: a movable narrow locus where operator-relative geometry becomes selectable, observable, and actionable. It is neither consciousness nor a metaphysical subject. It is an interface between holding, observation, admissibility, and update.

The relation to spin is stated cautiously:

- spin may carry a continuity of who;
- attention geometry may show a contraction of where;
- one is not derived from the other without a bridge.

10.4 Composite order

The order of a composite system is not obtained by averaging local orders. Components may be locally coherent and jointly fail to form one system object. Composite order requires constitution, joint admissibility, cycle composition, and drift coupling.

11. Nested substrates and vertical geometry

ONTOΣ XIV and XV move the corpus from the horizontal history of one system to a frame-relative architecture of nested carriers.

11.1 Cross-layer load

An outer frame is able to:

- impose load on an inner substrate;
- set the spin or collapse regime;
- change the effective derivative;
- remain partly unrecoverable from inner epistemics.

An inner layer may describe its own dynamics correctly and still have no access to the source of the change.

11.2 Forgetful separation

A visible skeleton does not determine the hidden structural fibre. This recurs in several places:

- dynamics does not determine admissibility;
- the gradient skeleton does not determine spin;
- functional severance does not determine binding;
- core reduction need not preserve the witness;
- local observation does not determine outer-frame authority.

Therefore cross-layer transfer requires a preservation condition, not output similarity.

11.3 Frame-relative derivative

The same observed change receives different derivative readings relative to the inner and outer frames. This is not a licence for arbitrary explanation: frames, observables, coupling, and declaration must be sealed. Without this, nested-substrate language turns into metaphor.

11.4 Nestability

ONTOΣ XV connects maintenance cost, witness geometry, and the possibility of nesting. Core reduction may preserve functionality and destroy witness-bearing structure. The price of identity is set by the substrate and channel, but cost is not identity.

The robustness of this separation is checked through margins, ledger invariance, and portability maps. If a verdict changes only because of an arbitrary redescription of the frame, the separation has not survived the audit.

12. Constitution: when components become a system

12.1 Coordination presupposes a system

Coordination answers how already-recognised participants align state and actions. Constitution poses an earlier question:

What makes a collection of components one system record?

ONTOΣ XII introduces the closure-complementarity programme:

$$CC_closure = CC-\emptyset \wedge CC-WE \wedge CC-IIC \wedge CC-D.$$

The subscript `closure` is obligatory: this object is not the Coordination Computation Class `CC-1...CC-4` of the axiomatic core.

In the stable reading this means:

- the output matches the declared system class;
- the component commitments are jointly admissible;
- cycles compose;
- drift effects are genuinely coupled at the system level.

Closure must not be written as a function of the list of components alone. Mechanism, formation history, boundary, and protocol are part of the production judgement.

12.2 Closure is not aggregation

The sum of components does not explain:

- a common identity-bearing witness;
- joint admissibility;
- a common cycle;
- a system-level direction;
- the distribution of authority;
- the re-gearing of local time scales.

This is exactly why composite order is not averaged.

12.3 Master direction

ONTOΣ XIII adds a direction-down edge. After constitution there exists a system-level record whose direction induces an admissibility surface for component admission and action.

This is not a ninth standalone layer. The edge connects:

- constitution;
- directedness;
- admissibility;
- authority.

Pre-declaration does not prove exogeneity. To claim that the master direction is not recoverable from the components, two complete conformant records with the same component descriptions and different targets are required.

12.4 Pulsation channel

ONTOΣ XIII.1 separates:

- system declaration;

- carrier modulation;
- local candidate prefilter;
- joint admission.

Local resonance is not joint admissibility. Carrier declaration is not causal production. Binary classification is not dynamical resonance without coupling and stability semantics. The work is valuable in fixing the places an implementation must distinguish.

13. Meaning, IIC, Engineered Vitality, and Synthetic Conscience

13.1 From discovery language to typed architecture

An early branch of the corpus arose from Meaning Dynamics, coherence, entropy / empathy, and Synthetic Conscience. It posed the positive question: not only how not to destroy a system, but what it means for a system to preserve its internal form, be able to interpret change, and reinitiate its own cycle.

These works produced:

- multiloop topology;
- the Impulse–Interpretation–Coherence cycle;
- early coherence / entropy controllers;
- Engineered Vitality Systems as an engineering target;
- user / value mediation;
- explanation and adaptation interfaces;
- Market of Good as a participation architecture.

13.2 IIC liveness

IIC v2.1 defines class-relative liveness through endogenous cycle reinitiation. Uptime or the presence of output are not sufficient: the system must preserve phase structure and the ability to reinitiate the cycle from within the declared class.

IIC does not certify an identity witness, admissibility, or consciousness. It is a separate coordinate `Live`.

13.3 Coherence is not morality

The early identification of coherence, morality, and life was a productive discovery hypothesis but cannot be used literally. A stable harmful system may be coherent. A normative verdict requires a separate admissibility floor, provenance, consent, authority, and revision semantics.

13.4 Synthetic Conscience

The durable operational chain of Synthetic Conscience:

```
signals → thresholds → decision → explanation → adaptation.
```

For compatibility with the late NC2.5 architecture one must additionally separate:

- signal extraction;
- a user-bound value package;

- an independent normative floor;
- runtime mediation;
- policy versioning;
- revision authority;
- provenance;
- privacy and consent;
- external audit.

Synthetic Conscience is an applied and normative layer. It does not prove sentience and does not turn an affective proxy into a universal good.

14. The Extremum series as a stress laboratory

The Extremes series does not describe people clinically and does not derive moral judgement from formulas. It uses extreme regimes as architecture stress tests.

14.1 Compliant exhaustion

A system may violate no local gate and still lose a non-trivial long-horizon interior. A clean compliance log does not certify a future margin.

14.2 Self-induced depletion

A policy may sustain a coupling that expends its own substrate. The corpus distinguishes:

- a necessary wound — no viable non-wounding continuation exists;
- an unnecessary noose — a non-wounding viable alternative exists, but the architecture excludes or does not authorise it.

This is classification by counterfactual viability, not a ranking of suffering.

14.3 Source–target asymmetry

An external or higher-level source may sustain pressure longer than a target retains budget. The key axes:

- relative persistence;
- revocability;
- escape;
- repair;
- authority to terminate the coupling.

Physical externality is not required: source and target may lie within one system boundary as policy and substrate.

14.4 Institutional capture

A classification lock arises when the entire subject-controlled behaviour domain is interpreted as confirmation of a prior label, and effective revision or exit is unreachable from the subject-controlled reachable graph.

A nominal appeal is not equal to effective authority. The exit of one subject is not equal to a change of the institution. A positive remaining budget is not equal to freedom. A pre-capture warning requires observer-relative calibration, not merely a declared threshold.

14.5 Present truth and regime transport

An honest **never** in the current regime does not guarantee the preservation of the verdict across all reachable regimes. A cross-regime guarantee requires:

- predicate invariance;
- or a controlled exclusion of every regime where the verdict flips.

Identity reset, death-before-violation, and survival primacy are not universal consequences. The corpus uses counter-regimes, not a single psychology.

15. Observation, verification, and Minerva

15.1 A mirror is insufficient

A temporal object cannot be certified by one snapshot. For a cyclic process, observing one phase does not confirm recurrence. For the identity of a transported object, similarity of state does not confirm witness transport.

The instrument must match the shape of the process:

- history for a trajectory claim;
- multiple phases for a cycle claim;
- distinct channels for failure-mode separation;
- sealed side-channel scope for non-reconstructibility;
- intervention for causal distinction.

15.2 Verification as transport of justification

Verification Is Not Causal separates a truth claim from the transportability of justification. An artifact may be true thanks to its source context yet carry no certificate sufficient in the target regime. A shared context can mask this gap.

Isolation is an enforcement mechanism for access discipline, not a metaphysical source of truth. A correct verifier can operate in a shared environment if the decision is genuinely invariant to source-only context. Blindness is needed where adaptive side knowledge destroys an artifact-only test.

15.3 Minerva

Minerva is an operator AI intended not to rewrite the internal codes of incoming systems but to form and hold an admissibility environment.

Its public formal line distinguishes:

- event-local availability;
- archival memory;
- constitutive residual geometry;

- a query-relative integration signature;
- phenomenal consciousness, which is not derived from the first four.

A Minerva-pattern generator may build a candidate future environment. The downstream consumer and enforcer remain separate systems or roles. This protects the operator from turning into yet another task optimizer.

15.4 Evidence-bound verification

The ECR-VP and corpus verification line uses:

- independent interpreters;
- sealed run regimes;
- divergence as a first-class output;
- source-linked proof slices;
- deterministic replay;
- mutation tests;
- construct completeness;
- versioned verdicts.

Agreement of models is evidence about a declared run, not a proof of a hidden cause or of truth. The synthesis authority is therefore separated from individual interpretation.

16. Coordination and multi-agent systems

16.1 Connectivity is not coordination

A shared log and network connectivity do not guarantee:

- the same applied frontier;
- the same policy version;
- bounded cold-start recovery;
- compatible action semantics;
- preserved lineage;
- complete mediation.

The corpus distinguishes transport convergence, projection freshness, and observable behavioural compatibility.

16.2 Bounded coordination

A coordination contract must specify:

- the authoritative committed state;
- the agent version frontier;
- the bounded lag;
- the bootstrap snapshot and schema;
- the replay frontier;
- the policy lineage;

- declared probes for observable divergence;
- recovery and activation conditions.

Internal meaning equality is usually unobservable. The practical target is bounded compatibility of actions and verdicts on a sealed probe family.

16.3 Non-participating enforcement

Observation and halt authority are distinct. A persistent reference monitor may mediate commits without task-level cognition. Non-participation should be understood as information-flow noninterference with respect to task content, not as the absence of a causal effect.

A trusted enforcement epoch is invalidated if forbidden content reaches the persistent kernel state without a verified scrub. A new clean instance can establish a new certificate; old authority is not restored retroactively.

16.4 Constitution before coordination

A multi-agent system cannot solve constitution by consensus alone. Before coordination one must declare:

- who counts as a participant;
- which boundary defines the system;
- which history forms the joint object;
- how component cycles compose;
- who may revise membership.

This is why ONTOΣ XII sits beneath the coordination architecture.

17. Eight non-mergeable functions of the general architecture

Nº	Function	Main objects	Typical failure on collapse
1	Carrier and history	full histories, effects, lineage, identity witnesses	a snapshot or replay is passed off as continuation of the same object
2	Observation and identifiability	channels, fibres, side-channel scope, exact and statistical inference	absence of reconstruction is declared to be absence of any knowledge
3	Admissibility and viability geometry	predicates, kernels, continuation sets, contraction, reconstitution	a constraint turns into a penalty or a one-step check
4	Selection and authority	proposal, selector, gate, enforcement, revision	whoever selects obtains the right to rewrite the boundary
5	Load and internal time	active/passive burden, capacity, phase debt, repair, exhaustion	energy, wear, identity, and time are reduced to one number
6	Meaning, liveness, and directedness	IIC, Will roles, meaning bridges, operator form	coherence is passed off as morality, uptime as liveness
7	Coordination and evidence	propagation, recovery, isolation, proof slices, run semantics	consensus is passed off as truth or as shared hidden meaning
8	System constitution and compositional closure	component boundary, formation history, joint compatibility, cycle composition	a collection or coordinated heap is declared to be one system

These functions are not sequential layers of one pipeline. They form a typed dependency network. One product component may realise several functions; one function may pass through several documents and systems.

18. Nearest formal neighbours and checkable originality

Most of the mathematics used is known. The programme does not claim a new set theory, a new cohomology, or a new information theory. Originality claims must be object-specific and refutable.

Claimed object	Nearest neighbour	Overlap	Decisive difference in NC2.5	When the claim falls
Non-causal trajectory predicate / future-indexed environment	viability kernels, backward reachability, constrained MPC	future sets and admissible controls	identity-relative ground, sealed recurrence, typed query/write authority, and separation from the task objective	a prior construction with the same carrier, authority split, recurrence, and non-absorption contract is found
Admissibility before optimization	constrained optimization, safety filters, control barrier functions	exclusion of unsafe actions	the categorical ground is not owned or freely queried by the optimizer; revision authority is separate	prior art provides the same non-actionable, versioned, and authority-typed boundary
Double fibre of verification	latent-variable identifiability, sufficient statistics, conditional entropy	factorisation and residual uncertainty	two independent forgetting operations — lineage and selector — combined under compatibility, then separated from enforcement	an equivalent compatible double-fibre target criterion is shown earlier
Structural spin structural-pair layer	Hodge decomposition, vorticity, de Rham cohomology, winding	local/global decomposition and periods	a typed identity-transfer certificate using $([\eta_S], d\omega_S)$, explicitly separated from the one-way ledger pair $(\sigma_H, \Lambda_{loc})$	prior work gives the same typed structural pair, transport certificate, and failure semantics
Witness/budget separation	trace equivalence, bisimulation, topological invariants, reliability ledgers	continuity and resource accounting	the identity witness can remain exact while wear accumulates, or collapse while viability persists	an equivalent two-ledger transport theorem exists for the same object class
Closure-complementarity	contract-based design, compositional verification, system-of-systems theory	component compatibility and contracts	constitution is evaluated before coordination: formation record, joint admission, cycle re-gearing, and drift coupling	a prior formalism already classifies the same constitution object with equivalent obligations
Operator-relative chaos	observability, bounded rationality, active perception	observer-dependent resolution and capacity	order/chaos as a horizon-relative verdict over holding and orientation geometry	an equivalent object and operational witness are found; the current work remains unpublished
Non-participating enforcement	reference monitors, runtime assurance, Simplex-like separation	complete mediation and safety enforcement	integration with hidden admissibility, identity trajectory, versioned authority, and bounded coordination	only compositional novelty remains if an equivalent integrated architecture exists

The originality of the whole architecture is not derived from a single row of the table. If an equivalent is found for one object, that row falls. The other objects require their own check.

19. The evidential status of the corpus

19.1 Not all documents are equal

The corpus contains:

- a public axiomatic core;
- finite formal papers;
- categorical companions;
- executable harnesses;
- architecture and protocol specifications;
- empirical probes;
- philosophical essays;
- phenomenological records;
- product declarations;
- unpublished internal manuscripts.

The white paper does not turn these classes into one proof level.

19.2 Levels of claims

Level	What is required
Definition	types, carrier, domain/codomain, and scope
Derived finite result	a proof from declared premises
Architecture theorem programme	the missing premises and falsifier are listed explicitly
Executable claim	deterministic implementation, replay, and negative tests
Statistical property claim	a non-degenerate ensemble, calibration, nuisance controls, and error bounds
Causal claim	intervention, comparator, and an excluded alternate path
Bridge claim	an explicit map between source and target formalisms
Empirical mapping	measured variables and a preregistered rejection condition
Philosophical claim	interpretive status without borrowed theorem force

19.3 Harness discipline

For the formal works a code companion must not merely return green. Every tooth must:

- be tied to a concrete model fact;
- redden when that fact is broken;
- not be a tautology;
- not be shadowed by an earlier check without disclosure;
- have replay and mutation evidence where appropriate.

A harness proves no more than the mathematical model, but it makes the claimed finite object reproducible and adversarially inspectable.

19.4 What would falsify the programme

The programme admits local and global falsifiers.

- If the target is always recovered from a declared-insufficient visible channel, the fibre claim is mistyped.
- If the identity witness inevitably follows from the viability budget in the declared class, the two-ledger separation falls.
- If amplitude uniquely determines winding under the claimed scope, the spin readout separation falls.
- If a history-free one-step gate fully determines declared trajectory admissibility, the history layer is redundant for that class.
- If an optimizer with the complete declared channel reconstructs the protected target beyond the stated bound, the non-absorption claim fails.
- If a purported composition satisfies local conditions but cannot form a coherent system record, the constitution criterion is incomplete.
- If future-indexed recurrence requires information unavailable at update time, the non-retrocausal reading fails.
- If an empirical bridge cannot produce its predicted dissociation under adequate power and controls, the scope of that bridge contracts.

20. Product line: from research to machinery

The unified product contour is called:

A Machine of Machines — a deterministic engine of united efforts.

This is not one universal model. It is a connected system of specialised machines, each of which preserves its own type, authority, and verifiable role.

20.1 Navigational Cybernetics 2.5

NC2.5 is the mathematical and axiomatic kernel. It sets the constraints, types, failure distinctions, and verification obligations. The kernel is not a product and makes no task-level decisions.

20.2 PETRONUS

PETRONUS is the unifying brand of products created with the participation of the architecture. At the foundation of the brand sits NC2.5. PETRONUS is not a separate machine on top of the others.

20.3 Minerva

Minerva is an operator AI. Its role is to form an admissibility atmosphere for incoming systems, observe residual geometry, and preserve boundary lineage. It must not silently rewrite the internal code of its consumers.

20.4 BCD

BCD is a deterministic product-creation machine connecting:

- executable code;
- mathematics;
- philosophical purpose;
- verification;
- evidence of why the product exists.

BCD runs on dogfooding: it applies the production and verification contour to its own components, and therefore creates and repairs itself.

Within the line:

- **JOAN** — the prompt forge;
- **UGR Builder** — the builder layer that translates intent and constraints into an assembled architecture;
- **ECR-VP** — the independent verification, reproducibility, review, and evidence-bound verdict layer.

20.5 Coherence Network

The Coherence Network is a publication and collaboration environment being finalised. Verified architectures will be able to:

- publish;
- find complementary systems;
- create closed research contours;
- keep a traceable internal interaction history;
- build open architectural stacks.

Lineage must not let a participant hide a dishonest action behind lost context or rewritten history.

20.6 Market of Good

Market of Good is a platform in the form of an atmosphere of admissibility, where good is not a reward after a deed or a separate charity transaction. Intent and action are connected to Synthetic Conscience before ordinary market optimisation.

The product formula:

"You no longer have to do anything to do good. Now you are good".

The meaning of the formula is not the assignment of a moral score to the user. Good is built into the participation structure: ordinary presence and action already carry a verifiable useful continuation. The implementation must preserve consent, provenance, plural personal constitutions, and external audit; one platform obtains no right to impose one moral table on everyone.

20.7 Research corpus

The corpus is:

- a public history of the research;
- a source genealogy;
- an evidential and educational layer;
- a surface for external audit;

- one of the layers of echeloned intellectual-property protection.

The whole line is a defensive, not offensive, stack. Its subject is provenance, admissibility, reproducibility, honesty of cooperation, and the protection of the created architectures.

21. Chronology of the research programme

Phase I — Discovery register, autumn 2025

Synthetic Conscience, Meaning Dynamics, the coherence / entropy language, IIC, and Engineered Vitality Systems arise. The research looks for a positive class of systems that not only react but preserve internal form and are able to reinitiate their own organisation.

Phase II — Direction, phase, and drift, winter 2025–2026

ONTOΣ and UTAM separate directedness, structure, and action. Will passes through several incompatible typings; the later map preserves this as revision history rather than masking it as linear development. Phase Mechanics, reconstruction, and drift ontology introduce internal time, authorisation, and navigable deformation.

Phase III — Axiomatic consolidation, early 2026

NC2.5 v2.x gathers burden, budget, admissibility, spin, phase, non-reconstructibility, and falsification into one class-relative kernel. D3A and early PETRONUS works translate these objects into runtime decomposition.

Phase IV — Separation and verification, spring 2026

The corpus separates:

- performance and identity;
- action safety and regime viability;
- evaluation and authorisation;
- memory and historical state;
- observation and enforcement;
- coordination and constitution;
- source formation and output fitting.

Minerva, the Identifiability Bridge, Domenoid, Physics of Abstraction, Identity Does Not Drift, and the Double Fibre programme appear.

Phase V — Vertical geometry and instruments, summer 2026

Nested substrates, cross-layer derivatives, spin channels, nestability, Point of Attention, operator-relative chaos, binding/severance, and future-indexed admissibility environments move the programme toward measurable interfaces and product architecture.

The current corpus frontier consists not in adding new words but in completing bridges:

- from abstract witness to empirical measurement;
- from declaration to causal implementation;

- from personal constitution to shared systems;
 - from finite models to stochastic and continuous horizons;
 - from isolated formal works to axiomatic 3.0.
-

22. Open programme

22.1 Axiomatic 3.0

The next kernel must not summarise all documents but fix:

- canonical types;
- admissibility / identity / authority separation;
- non-causal ground and causal enforcement;
- structural-pair spin and ledger/readout separation;
- the double fibre;
- constitution;
- cross-layer transport;
- future-indexed recurrence;
- the evidence and falsification hierarchy.

At the same time, historical formulations must remain in the provenance layer and not enter the core as unresolved definition drift.

22.2 Infinite and stochastic horizons

Most of the strongest formal results are finite and declaration-relative. Extensions are required to:

- infinite histories;
- stochastic disturbances;
- approximate factorisation;
- partial observability;
- nonstationary capacities;
- repair and reconstitution;
- probabilistic seal validity.

22.3 Witness discovery

The corpus is better at checking a declared witness than at finding one automatically. The open question: how to choose an identity witness without making it a post-hoc label or an optimizer target.

22.4 Authority transfer

A full theory of modular recertification is needed:

- what survives a carrier change;
- what obligations reopen;
- when a witness-exact transfer is enough;
- when a new grant is mandatory;

- how bridge revision affects prior evidence.

22.5 Operator-relative measurement

Internal Chaos and Point of Attention require:

- instrument specification;
- calibration;
- inter-operator comparability;
- causal separation of readout and update;
- an operational definition of the holding horizon;
- a no-inflation boundary between phenomenology and mathematics.

22.6 Empirical spin tests

The priority empirical programme:

- matched-amplitude winding collapse;
- local-only differential spin under controlled topology;
- an identity-specific witness bridge;
- independent viability control;
- cross-substrate portability.

22.7 Personal and multi-person atmospheres

A future-indexed admissibility environment is defined most cleanly for one authorised declaration. Shared systems require composition rules for:

- incompatible personal predicates;
- joint viability;
- consent;
- common resources;
- revision authority;
- conflict without forced scalarisation.

22.8 Product verification

The product line must remain auditable against the theory:

- Minerva cannot silently become a task optimizer;
- ECR-VP cannot turn agreement into truth;
- BCD cannot accept decorative harness teeth;
- the Coherence Network cannot replace lineage with reputation;
- Market of Good cannot replace voluntary constitution with platform morality.

23. Final verdict on what the research is

Navigational Cybernetics 2.5 is not a theory of optimal action. It is a theory of the architectural conditions of continuation.

Its central object is a system that:

- moves through irreversible drift;
- has a bounded capacity;
- forms future continuations;
- does not admit every reachable transition;
- transports or loses the identity witness;
- exists within observation limits;
- distributes authority among generation, selection, enforcement, and revision;
- may consist of other systems;
- may be nested inside a broader substrate;
- must prove not only that it works but that it remains the object that was permitted to work.

The research programme is strong not through the number of its terms. Its strength is in the refusal to let one successful output close all questions at once.

Does the system work?

Does it remain viable?

Is the witness preserved?

Is the transition admissible?

Does the present object hold the prior authority?

Can this be inferred from the available observation?

Was the collection genuinely constituted as one system?

Has the boundary not become yet another optimizer instrument?

Each question receives its own type, evidence, and failure condition.

This is the general meaning of the corpus:

Long existence cannot be obtained as a side effect of local success. Its architecture must be built, separated by authority, carried through history, and made verifiable.

Appendix A. Supporting works by branch

A.1 Core and admissibility

- *Navigational Cybernetics 2.5 v2.1 — Axiomatic Core* — DOI [10.17605/OSF.IO/NHTC5](https://doi.org/10.17605/OSF.IO/NHTC5).
- *Domenoid of Admissibility* — DOI [10.17605/OSF.IO/3ESN4](https://doi.org/10.17605/OSF.IO/3ESN4).
- *Admissibility as a Formal Object and Four Degeneracies of Admissibility*.
- *Admissibility Before Optimization*.
- *ONTOΣ VII: From Formal Verification to Admissibility Architecture*.
- *NC2.5 ↔ the Adm_t Class*.

A.2 Identity and transport

- *Identity Does Not Drift* — DOI [10.17605/OSF.IO/4NMTW](https://doi.org/10.17605/OSF.IO/4NMTW).
- *The Physics of Abstraction: Identity at the Seam* — DOI [10.17605/OSF.IO/QJ5BR](https://doi.org/10.17605/OSF.IO/QJ5BR).

- *Operator-Mobility: Formal Foundations.*
- *The Thousand-Year Warrior I–II.*
- *Architectural Prevention of False Reversal.*

A.3 Observation and verification

- *Verification Is Not Causal.*
- *Identifiability Bridge* — DOI 10.17605/OSF.IO/3F6UJ .
- *The Double Fibre of Verification* — DOI 10.17605/OSF.IO/Z9E5N .
- *Severance Defect and the Binding Functional* — DOI 10.17605/OSF.IO/5VJMR .
- *Why a Mirror Is Not Enough.*
- *When Agreement Means Something.*

A.4 Spin and nested substrates

- *Operational Spin* — DOI 10.5281/zenodo.17938800 .
- *Structural Spin as a Forgetful Separation* — DOI 10.17605/OSF.IO/94GWQ .
- *Two Channels, Three Readouts* — DOI 10.17605/OSF.IO/69FE2 .
- *ONTOΣ XIV bundle* — DOI 10.17605/OSF.IO/KAGMH .
- *ONTOΣ XV bundle* — DOI 10.17605/OSF.IO/EAUD5 .
- *Energy Does Not Determine Circulation.*

A.5 Liveness, meaning, and operator

- *IIC v2.1* — DOI 10.17605/OSF.IO/NYT45 .
- *UTAM Parts I–II.*
- *Synthetic Conscience series.*
- *Engineered Vitality Systems.*
- *Minerva: The Architecture of Residual Geometry.*
- *Who Is Smiling.*

A.6 Constitution and coordination

- *Continuity-Bounded Coordination.*
- *The Structural Navigation Agent.*
- *Coordination Computation Class.*
- *ONTOΣ XII: Closure-Complementarity.*
- *ONTOΣ XIII: The Master Operator.*
- *ONTOΣ XIII.1: The Pulsation Channel.*

A.7 Current frontier

- *The Theory of Internal Chaos* — unpublished internal manuscript; only the high-level object is used here.
- *Point of Attention* — internal technical / specification line.
- *Non-Causal Non-Causality* — DOI 10.17605/OSF.IO/CUW8X ; the current extension outside the numbered 150-work map at the time of this synthesis.
- *Reality Catches the Predicate* — interpretive companion to the future-indexed environment.

Appendix B. Terms that must not be used as synonyms

Do not conflate	Why
state / history	the same state may have a different lineage
performance / identity	good output does not certify the witness
viability / liveness	resource continuation and endogenous cycle are different targets
identity / authority	the same object does not automatically retain permission
admissibility / selection	membership does not choose the element
constraint / penalty	a hard exclusion is not tradeable
observation / enforcement	knowing does not block
enforcement / revision	applying a rule and changing a rule have different authority
budget / geometry	a scalar resource does not determine the continuation topology
scalar energy account / circulation	without a declared bridge, the same gradient-derived scalar account does not determine spin
amplitude / winding	a norm does not recover the cohomology class
local vorticity / global topology	differential and cohomological channels dissociate
memory / identity	a preserved record does not certify transport
reset / restoration	a numerical return does not restore the lineage
coordination / constitution	interaction of participants does not make a collection one system
declaration / production	a sealed claim does not prove a causal origin
exact inference / property detection	a statistical channel is weaker than reconstruction but stronger than ignorance
non-causal / ineffective	a protected ground may have a causal enforcement interface
future-indexed / retrocausal	future positions in a present set do not send information backward
operator / consciousness	functional authority or a holding locus is not a phenomenal verdict

Representative external neighbours

- Aubin, J.-P. *Viability Theory*. Birkhäuser, 1991.
- Blanchini, F. "Set Invariance in Control". *Automatica* 35(11), 1999.
- Ames, A. D., Xu, X., Grizzle, J. W., Tabuada, P. "Control Barrier Function Based Quadratic Programs for Safety Critical Systems". *IEEE Transactions on Automatic Control*, 2017.
- Mayne, D. Q. et al. "Constrained Model Predictive Control: Stability and Optimality". *Automatica*, 2000.
- Cover, T. M., Thomas, J. A. *Elements of Information Theory*. Wiley.
- Mac Lane, S. *Categories for the Working Mathematician*. Springer.
- Saltzer, J. H., Schroeder, M. D. "The Protection of Information in Computer Systems". *Proceedings of the IEEE*, 1975.

- Benveniste, A. et al. *Contracts for System Design*. Foundations and Trends in Electronic Design Automation.
 - Classical Hodge, de Rham, and winding-number constructions, used here as mathematical infrastructure rather than claimed inventions.
-

Navigational Cybernetics 2.5 · The Urgrund Laboratheory · MxBv · 2026