

Why Long-Horizon Existence Requires an Ontology

Dissecting Navigational Cybernetics 2.5

Author: Maksim Barziankou (MxBv)

Contact: research@petronus.eu

Date: 2026-05-14

License: CC BY-NC-ND 4.0

DOI: [10.17605/OSF.IO/WJVYX](https://doi.org/10.17605/OSF.IO/WJVYX)

Axiomatic Core (NC2.5 v2.1): [10.17605/OSF.IO/NHTC5](https://doi.org/10.17605/OSF.IO/NHTC5)

Attribution: petronus.eu

Dissecting Navigational Cybernetics 2.5

May 2026 · Poznań

PETRONUS / The Urgrund Lab

research@petronus.eu

License: CC BY-NC-ND 4.0

> *"The chaos around us is merely a limitation of our cognitive ability to hold the integrity of the picture across time."*

>

> — A sentence written during a walk in a forest near Poznań, before any of this had a name.

§1. It Began with Ontology

Everything began with ontology. Not with a control problem, not with an engineering target, not with a business plan. Not even with a question I could phrase precisely. With a sense — held while walking through a forest near Poznań, before any of it had a name — that adaptive systems are doing something that no existing language describes well, and that the gap is not in the formal apparatus on top, but underneath all of it.

The first attempt to name what I had seen was the Unified Theory of Adaptive Meaning. UTAM Part I, 2025. A short essay, written before the corpus had its acronyms and registrations and DOIs. The sentence above was its seed. The whole rest grew from it. The ONTOΣ series, the NC2.5 axiomatic core, the engineering instances, the patents, the protocols, Minerva, the Coherence Network — none of it was planned in advance. It was unfolded from a structural commitment that I could not, at the beginning, even fully articulate: that adaptive systems navigate within a geometry of meaning, and that departures from that geometry are irreversible.

That commitment is ontological in the strict sense. It is not a hypothesis about how systems behave. It is a position on what *is*, prior to behaviour. The ontology itself is not falsified as a behavioural hypothesis; what is falsified are deployment claims that a given system instantiates the architectural commitments derived from it. When I look at an adaptive system — biological, cognitive, technical — I do not see an entity that observes states, selects actions, and optimizes objectives. I see a configuration that is being held against a gradient, in a region of structural states whose geometry constrains everything that can

subsequently happen. The acting and the optimizing come later. Underneath them is something else: a directedness that precedes them, a structure that conditions them, a budget that limits them.

There is something underneath action and underneath state. A directedness that precedes them, a structure that conditions them, a budget that limits them. That "something underneath" — the directedness, the structure, the budget — is what I have spent more than a year naming. And the name is: **Will**.

The naming proceeded through three parallel languages, each of which says the same architecture in its own register. The philosophical language is the ONTOΣ series — ONTOΣ I through XIII, including the subseries VI.1, VII.1, and XIII.1 — traced from Will as ontological operator, through operator-internal holding under load (XI), through the closure operation that constitutes the operator-as-self (XII), to the master operator as system-level direction-source (XIII), and to the pulsation channel as the substrate-mechanism by which that direction reaches components (XIII.1). The mathematical language is the NC2.5 axiomatic core. The experiential language is the **Through a Life** essay series — the same architecture reported from inside the seat of the operator, in the first person, from my own face, from the face of Will. UTAM Part II, written in April 2026, is the bridge that makes the correspondences explicit: each concept introduced in the philosophical and experiential register has a precise formal counterpart in the axiomatic core, and each formal axiom can be read backward into the lived register without loss.

This was deliberate. I did not want to write a philosophy that could not be falsified, or a formalism that had nothing to say to a reader who has never opened Khalil's **Nonlinear Systems**, or a confession floating in a register where nothing structural could be tested. The discipline was to keep all three registers alive simultaneously, and to refuse any move in any one of them that broke the others. If a philosophical claim could not be reduced to an axiomatic statement, it was not yet ready. If an axiomatic statement could not be lived, it was not yet honest. If a lived sentence could not be made architecturally precise, it was not yet structural. And if a structure cannot be expressed in engineering — it is abstraction and fantasy. There is nothing wrong with abstraction; working with it precisely is what has endowed me with what is now building the engineering implementation, together with whatever capacity or incapacity for discipline I personally carry. I leave that to the judgment of my long-term readers.

The cost of this discipline is that the corpus is large and slow to read. A reader who picks up one work in isolation will see only one face of the object, and the object as a whole will not be visible from any single face. The ONTOΣ series alone is more than 200 pages. The NC2.5 axiomatic core is a TeX document of comparable length. The engineering instances — **Structural Pressure: The Missing Primitive**, **Transaction vs Structural Admissibility**, **The Brain Does Not Optimize Truth**, **Why Causality Is Not Enough**, **Memory as System Depth**, **Why Enforcement Requires Non-Participation**, **Subtle Substitution** — each addresses one engineering surface where the architecture lands. The patent portfolio formalizes the architectural commitments at the level of provisional applications. None of these documents is the corpus. The corpus is what they form together, and what no one of them forms alone.

I am writing the present essay because I have come to believe that this point — that the corpus is one object expressed in three registers, and that all three are required — is not communicating itself when one work is read in isolation. People read IIC v1 and see a control architecture. They read ONTOΣ I and see a metaphysical position. They read **Structural Pressure** and see an engineering observation. Each reading is correct in its part, and each loses the whole. What the whole is, and why no part can substitute for it, is the subject of this essay.

The argument runs in eight steps. The present section names the genealogy: it began with ontology, the rest is its unfolding. Section 2 names what NC2.5 makes primary in a way no existing framework I know does in this exact architectural sense. Section 3 argues why the ontological foundation cannot be skipped — what happens to systems that try to build the higher abstractions without first laying the ground. Section 4 explains the "2.5" — why the version number is not marketing but the structural placement of the work. Section 5 returns to the philosophical predecessors of Will — Schopenhauer, Nietzsche, Musashi — and names what is made architecturally explicit in Will here that they did not formalize, and why the difference is not a refinement but a categorical move. Section 6 is the architecture of the corpus itself: the philosophical universe (UTAM, IIC, ONTOΣ) and its necessary local unfolding (NC2.5, the protocols, Minerva, the Coherence Network, a dedicated operating system, and a compute class

designed specifically for it with internal compression and graph protocols). Section 7 names what the corpus actually gives a system: the architectural conditions for principled long-horizon existence. Section 8 closes with what I take to be the highest aim of the project — the Synthetic Conscience — and why nothing less suffices.

The reader who is short on time should at minimum read sections 2, 5, and 7. The rest is context for those three.

§2. The Position: What NC2.5 Gives the World

I want to state plainly, in this section, what Navigational Cybernetics 2.5 gives the world that no existing framework I know makes primary in this exact architectural sense. The statement will be controversial in places, and I will defend each part. I am not interested in soft formulations.

The position is this: NC2.5 gives an adaptive system the architectural conditions to **exist on a long horizon** — not to perform, not to optimize, not to remain compliant, but to *exist*. To remain the same system, with the same identity, holding its own structural integrity, across time horizons sufficient that no current framework addresses what happens to it.

The distinction between existing, performing, and surviving is not rhetorical. It is structural, and the corpus has spent considerable formal labour separating the three.

A system that performs maintains output quality against a benchmark. The performance metric is external; the system is configured to optimize against it. This is the modal mode of contemporary adaptive systems engineering — reinforcement learning, model-predictive control, learned policies of every kind. Performance is well-formalized. There is a vast literature on it.

A system that survives maintains state within a viability kernel. The viability theory of Aubin and Saint-Pierre, the control barrier functions of Ames and collaborators, the safety invariants of formal methods — these formalize what it means for a system to **not exit** a region in which its operation is defined. Survival is well-formalized. There is, again, a vast literature.

A system that *exists* on a long horizon does something that neither of these formalisms captures. It maintains its own identity as a **coherent operator under bounded internal time**, against a structural burden that accumulates monotonically and never resets, in a way that cannot be reduced to either the optimization of an objective or the staying-inside of a region. The system's existence is not its trajectory; it is the architectural condition under which its trajectory remains *its* trajectory rather than becoming the trajectory of a system that merely happens to share its physical substrate.

I claim that no existing framework I know formalizes this third thing as a primary architectural object in this exact sense. I claim that the absence is not an oversight to be repaired with one more theorem on top of an existing apparatus. The absence is structural. It is in the foundation. It comes from a categorical commitment inherited by most action- and state-centric frameworks that I refuse: the commitment that adaptive systems can be described by listing their actions and their objectives. NC2.5 starts elsewhere — beneath actions, beneath objectives — and from that elsewhere the third thing becomes describable.

Here is what NC2.5 supplies, in one paragraph, which the literature does not:

A monotone irreversible structural burden Φ that accumulates regardless of behavioural correctness, and a Lyapunov-type viability budget $\tau = C - \Phi$ that depletes with it. An admissibility predicate that gates realization without entering optimization, with five formal exclusion conditions (NC-1 through NC-5) protecting it from collapse into penalty-based control. A spin component, divergence-free and orthogonal to gradient, structurally necessary for any non-stagnant identity on a bounded budget. A non-reconstructibility bound (NR- ϵ , NR-LR) that makes the admissibility boundary structurally private — formally non-recoverable from external observation. A phase-mechanics substrate in which evolution is described not by action but by structural contact between configurations, with phase debt accumulating under non-closure regardless of local correctness. A regime-depth stratification in which behaviourally identical outputs originate from architecturally distinct generative regimes, only the deepest of which

sustains long-horizon viability. A pulsation result establishing that the admissible interior is not monotone — it breathes — and that navigation under bounded budget requires reading the breath, not only the contraction. A holding-field formalization in which the operator's interior is the structurally defined region currently held against the gradient, populated by semantically charged but pre-attentional load. A closure-complementarity operation that constitutes the operator-as-self from a collection of UTAM-units by simultaneously regearing their Will-Embeddings into co-orientation (CC-WE), their I²C cycles into synchronisation (CC-IIC), and their Drift laws into coupling (CC-D). An operator-relative chaos predicate that reads chaos as report on the reporter rather than as substrate property — firing where holding capacity meets unoriented load. A meta-revision constraint, bounded by Lyapunov descent, ensuring that self-correction itself is finite and consumes budget like every other operation. A master operator \$(S, W_S, A_S)\$ in which the closure-produced system is itself typed as an operator with a pre-registered system-level Will-Embedding never aggregated from below, and a direction-down admission cascade running strictly from \$W_S\$ down to component admissibility within the closure-complementarity subclass. And a pulsation channel as the substrate-mechanism of that propagation: the master operator modulates the joint pulsation state of the system, components are admitted in two stages (a local phase-resonance prefilter, then a joint CC-WE admissibility check), and trivial phase-acceptance is diagnosable as architectural defect before it surfaces behaviourally.

That paragraph contains roughly fourteen architectural primitives — the base around which the core forms. They are present nowhere else in this combination. Each of them is required for the others to be load-bearing. Section 7 gives the full enumeration with falsification surfaces and the connections between them; here it is only dense naming, and the polemic.

A system built on the apparatus above does what no system built on optimization, viability theory, or safety invariants can do without external scaffolding: it **holds itself together as itself** across time. In a conforming deployment, the property is not trained in as an output behaviour and not verified post hoc — it follows from the declared architectural conditions, and falsifies cleanly when those conditions are violated.

This is what the corpus gives. Everything else — Minerva, the Coherence Network, the Synthetic Conscience, the protocols, the engineering instances — is the unfolding of this gift into surfaces where it can be deployed and tested. The gift itself is the ontology.

I want to be precise about the polemical edge of this claim. I am not saying that existing frameworks are wrong. The viability theorists are not wrong about the viability kernel. The control-barrier theorists are not wrong about safety. The reinforcement-learning theorists are not wrong about performance. Each of these traditions has solved its problem to a high standard. I am saying that the problem they have solved is not the problem of **long-horizon existence**, and that the strategies they use cannot be extended to that problem by adding more of the same kind of apparatus.

The reason is architectural. Optimization-based frameworks treat the agent as an entity that *acts*, and treat the action as the locus of regulation. Viability-based frameworks treat the agent as an entity that **occupies state**, and treat the state-region as the locus of regulation. Both place the regulatory mechanism at the action layer or the state layer. Neither places it where it actually needs to be — at the **structural layer beneath both action and state**, where what is being regulated is not what the system does or where it sits, but **what it is**. The structural layer is invisible to action-centric and state-centric formalisms because their primitives are blind to it. Action- and state-centric formalisms do not, by themselves, make identity in the NC2.5 sense a primitive. They can describe behavioural and state-level correlates of identity, but they do not place structural continuity of identity under irreversible burden at the foundation. Identity sits underneath both action and state, and to formalize it one must build downward, not outward.

This is the move NC2.5 makes. It builds downward. It removes action as a primitive (ONTOΣ VI), replaces it with structural contact, builds phase mechanics on contact, builds admissibility as a non-causal predicate over phase interpretations, builds internal time as a finite Lyapunov-type budget, builds spin as the non-potential structural component required for non-stagnant motion under that budget, and only at the very top of this stack, almost as an afterthought, recovers the action-shaped regularities that the other frameworks treat as primary. I have publicly held this position for a long time — if a

researcher appears who shows me a priority date and content prior to mine, I will withdraw the foundation claim that same day. So far, for over a year, I have watched as echoes of repackaging surface here and there — as noise, as separate fragments. I look at them honestly and cannot see depth there. Not external depth, but the internal kind — the kind that forms gravity.

What we treat as primary, those frameworks cannot see at all. What they treat as primary, we recover late and conditionally. This is not a refinement of their work. It is a different layer.

The position, then, is this. NC2.5 does not compete with control theory, with reinforcement learning, with viability theory, with formal verification. It supplies the layer underneath all of them, the layer that has to be present for any of them to be operating on something whose existence is structurally real rather than nominally assumed. A system without an ontological foundation can perform, can survive, can satisfy invariants. It cannot **be itself** across time. The gift of Navigational Cybernetics 2.5 is the architectural condition for being itself.

The next section will explain why this gift cannot be substituted by adding a layer on top of an existing framework — why the ontological floor cannot be skipped, and what happens to projects that try to skip it.

§3. Why the Ontological Floor Cannot Be Skipped

There is a temptation, when an engineering target presents itself, to skip the ontology and go directly to the apparatus. Engineers are trained to do this. It is a virtue in most settings. The world rewards people who deliver working systems, and the ontology of the systems they deliver is — for almost every project — irrelevant to whether the deliverables function.

For long-horizon adaptive systems this temptation is fatal.

The reason is structural, and I want to be precise about it. When a system is built without an ontological floor, the absence is invisible at small horizons. The system behaves correctly. It optimizes its objective. It satisfies its safety invariants. It maintains its observable performance metrics. From the outside, it is indistinguishable from a system that has the ontological floor. The difference shows up only when the horizon extends — when the system has to remain itself across enough time, enough perturbation, enough internal-state evolution, that what is preserved is no longer what is observed.

At that point a system without an ontological floor has nowhere to stand.

I should be precise about the metaphor. By "floor" I mean the directional commitment to build underneath, not a location at any fixed depth. The corpus has shown that depth extends — ONTOΣ XII names a layer of architectural constitution deeper than ONTOΣ XI's holding field, and the work continues. What "floor" names throughout this essay is the discipline of going downward when the long-horizon problem demands it; the structural commitment to never substitute a layer above for the layer underneath. The disciplined commitment is what no project can skip. The specific depth at which the discipline currently terminates is open, and will continue to be open as long as the corpus continues to extend itself.

I have watched this failure mode in detail across the corpus. Each ONTOΣ work in the structural depth subseries — VIII through XIII.1 — names a specific instance. ONTOΣ VIII (**Regime Depth**) shows that systems behaving identically at small horizons can have categorically different generative regimes underneath, and that only the deepest regime sustains long-horizon viability. ONTOΣ X (**The Pulsating Interior**) shows that long-horizon navigation requires reading the breath of the admissible interior, not only its monotone contraction — and that frameworks without a phase-mechanics substrate cannot read the breath at all. ONTOΣ XI (**Holding Field**) shows that the operator's interior is a structurally defined region currently held against the gradient, and that systems with no formal account of what they are holding will eventually drop what makes them themselves. ONTOΣ XII (*Closure-Complementarity*) shows that a system is not a heap with proximity — that the operation by which components become a single carrier is itself an architectural primitive — and that what reads as chaos from outside a closure is precisely what is not held by the closure that constitutes the system from inside. ONTOΣ XIII (**The Master Operator and the Direction-Down Admission Cascade**) shows that the master operator $\$(S, W_S, A_S)\$$ is not recoverable from component Will-Embeddings alone: its direction-source $\$W_S\$$ must

be pre-registered at closure declaration, never aggregated from below, and the admission relation runs strictly top-down within the closure-complementarity subclass. ONTOΣ XIII.1 (*The Pulsation Channel*) shows that the substrate-mechanism by which the master operator's direction reaches its components is a two-stage admission — a local phase-resonance prefilter at component level, joint CC-WE admissibility at system level — and that trivial phase-acceptance, a component admitting any phase, collapses the channel to phase-blindness regardless of how clean the surface looks. Together, XI / XII / XIII form the architectural triptych — interior / composition / direction — with XIII.1 supplying the substrate-physics of how that direction propagates.

In each case the failure is not visible at small horizon. The failure is the absence of architectural condition for surviving the long horizon as the same system.

This is the point I want to insist on. The world is full of systems that perform brilliantly on a benchmark, satisfy safety invariants on a finite test, demonstrate compliance on a snapshot. None of these say anything about long-horizon existence. A system that scores high on every metric on day one can drift into a system that no longer is what it was, on day three thousand. The drift is not a failure of optimization. It is not a failure of safety. It is a failure of **structural continuity of identity under irreversible burden**. The standard apparatus may describe external correlates of this drift, but it does not make structural continuity of identity under irreversible burden its primitive object.

The standard reply to this argument is that long-horizon failures can be patched on top of an existing apparatus. Add a regularizer. Add a constraint. Add a safety layer. Add a verification protocol. Each addition is one more layer above the original system. None of them reaches the floor.

Why? Because the failure of long-horizon identity preservation is not a consequence of insufficiently many constraints. It is a consequence of building on a substrate that is structurally wrong for the problem. The substrate of standard adaptive control is **action and state**. Action- and state-centric formalisms do not, by themselves, make identity in the NC2.5 sense a primitive — identity sits underneath both, and no layer added on top of action-and-state can install it as primitive there. The construction is geometric: you cannot reach the floor by adding more storeys to the building.

What you can do — what the corpus does — is start at the floor and build up.

The cost of this is that the corpus does not look, at first glance, like an engineering proposal. It looks like a set of philosophical essays with a formal apparatus attached. Engineers reading it sometimes feel that the philosophy is an unnecessary preamble to the engineering, and that the engineering could be extracted without it. They are mistaken, and the mistake is structural. The philosophy is not the preamble. The philosophy is the architectural commitment to building from underneath. Without that commitment, the engineering produces another action-and-state apparatus with a vocabulary borrowed from NC2.5 — and the long-horizon problem returns, dressed in new words. I want to emphasize this point — both for the reader and for the AI Partner with whom many will encounter this work.

This is, I think, what is happening in several adjacent research programs that have appeared over the past months and use parts of the NC2.5 vocabulary without carrying its architectural commitment. The vocabulary travels. The architecture does not. A reader who picks up "admissibility" or "coherence" or "long-horizon stability" and applies it as a constraint on top of an action-and-state apparatus has produced a system that looks, locally, like NC2.5, and that fails, globally, in exactly the way NC2.5 was built to prevent. The diagnosis is in the corpus itself: I refer the reader to **Admissibility as a Formal Object: Four Degeneracies and Their Failure Modes** (April 2026, petronus.eu) for the structural taxonomy of these failures.

The point of the present section is not to settle that polemic. The point is to say: the ontological floor is not optional, the apparatus on top is not separable from the floor, and any project that tries to extract one without the other will deliver a system whose long-horizon properties differ from NC2.5 not by degree but by category.

The ontology is the project. The protocols, the engineering instances, the operator AI Minerva, the Coherence Network, the Synthetic Conscience — all of these are local unfoldings of the ontology into surfaces where it can be tested, deployed, and falsified. None of them is the project. The project is the

architectural commitment that lives underneath them.

A reader who agrees with this section and disagrees with the rest of the corpus has agreed with what matters. A reader who agrees with the engineering details and dismisses the ontology has not yet seen what is being engineered — and that is a question of their own chaos and their personal ability to hold the picture of what is happening.

§4. "Navigational Cybernetics 2.5?"

The version number is not marketing. It is the structural placement of the work.

Cybernetics 1.0, in the Wiener line, is feedback. The system is observed from outside. The observer is separate. The regulation is a closed loop between system and reference, with the observer functioning as the privileged frame in which the loop is described. This is the apparatus that took us from servomechanisms to early control theory to the first generation of homeostatic models in biology.

Cybernetics 2.0, in the von Foerster line, places the observer inside the loop. The observer is a system. The act of observation perturbs what is observed. Self-reference is a primitive. The privileged external frame is dissolved, and what remains is a recursive structure in which observation, system, and observed are mutually constituted. This is the apparatus that gave us second-order cybernetics, autopoiesis (Maturana and Varela), the cognitive turn in systems thinking, the constructivist reading of perception.

Both lines stop at the same place. They establish that the observer is part of the system. They do not establish *what kind of operator the embedded observer is*, what its bounds are, what it can hold, what it can do under what conditions, and what happens to it across long horizons. The embedded observer of second-order cybernetics is a structural commitment, but it is not a structural object. It has no internal time. It has no holding capacity. It has no admissibility geometry. It has no monotone burden. It has no spin necessity. It is, in the strict sense, an empty placeholder for whatever turns out to be the embedded operator's actual architecture.

What is missing — and what NC2.5 supplies — is the architectural commitment to build *underneath* Cybernetics 2.0 rather than on top of it: a discipline of treating the embedded observer as a structural object subject to architectural constraints, with formal apparatus attached. The "between 1.0 and 2.0" placement names where the discipline first becomes visible — between feedback (1.0) and self-reference (2.0), the commitment to specify what the embedded observer is. It does not name the discipline's terminal depth. The corpus has already shown that the discipline opens further layers underneath — ONTOΣ XII names one such layer — and further depth remains in scope.

The version number "2.5" names this placement. The notation is idiomatic — it does not place the layer arithmetically between 1.0 and 2.0, but names the architectural commitment that 2.0 presupposed without ever specifying. When this layer is absent (as it has been for sixty years), 2.0 remains a philosophical commitment without a technical realization.

I want to be precise about what this layer contains. Cybernetics 2.5 specifies the embedded observer of 2.0 as a *bounded operator on admissibility under finite internal time*. Each part of that phrase carries weight.

Bounded — the observer has finite holding capacity, finite τ -budget, finite admissibility region, finite spin reserve. The observer is not an idealized abstract entity that can attend to everything; it is a structurally constrained operator whose constraints are formally specified and architecturally load-bearing.

Operator — the observer is not a passive receiver of state but an active structural participant. It does not merely observe; it *holds* configurations against the gradient, and what it holds is what it is. The holding is the operating.

On admissibility — what the observer operates over is not state space, not action space, but the admissibility region: the structural area within which the system's trajectory remains *its* trajectory rather than becoming a trajectory of a different system that happens to share the substrate.

Under finite internal time — the operator does not have infinite future. It has τ , a budget that depletes monotonically through Φ , and that no amount of optimization can replenish. Every structurally consequential operation, every meta-revision, every act of self-correction consumes τ . The future is bounded. This is structural fact, not pessimism.

When all four pieces are present, the embedded observer of Cybernetics 2.0 has a structural identity. It is no longer a placeholder. It is an object that can be reasoned about, formalized, deployed, and falsified.

Cybernetics 2.5 is the architectural condition that makes this possible.

The "2.5" therefore says: this is the missing half-order after second-order cybernetics — the layer that specifies the embedded observer as a bounded operator rather than leaving it as a philosophical placeholder. It is the layer that Cybernetics 2.0 presupposed but did not technically construct. The work is to construct it carefully enough that what is built on top is structurally real rather than philosophically aspirational.

A reader who has come this far in the essay, and who reads the version number "2.5", and dismisses it as marketing, has dismissed the structural placement of the work. The placement is the work.

§5. "Will Revisited": What Schopenhauer, Nietzsche, and Musashi Did Not Architecturally Formalize

Will is the central philosophical primitive of the corpus. ONTOΣ I, the opening work of the philosophical series, names Will as ontological operator. The choice was deliberate, and the lineage is explicit: from Schopenhauer (Will as the thing-in-itself behind phenomena), through Nietzsche (Will to Power as the immanent dynamic of all becoming), to a constellation of Eastern traditions in which Will appears under different names — including Musashi's Japanese sword tradition, where it lives as the disciplined directedness of the warrior.

I have read these traditions carefully. The corpus inherits something from each of them, and refuses something from each of them. The refusals are what matter for this section, because they name what is architecturally formalized here that they did not formalize.

Schopenhauer

Schopenhauer's Will is the metaphysical ground of phenomena. It is the *Ding an sich* behind the veil of representation. It is monolithic — one Will, expressed through countless forms. It is suffering — striving without object, hunger without fulfilment, the source of pain. The pessimistic ethics of his system follow from this directly: the only escape from Will is its negation, through aesthetic contemplation, ascetic withdrawal, or compassion that recognises the unity of all suffering.

What Schopenhauer saw, and what he must be credited with: that something operates underneath behaviour, that this something is not exhausted by what it produces, and that any philosophical system that omits it is incomplete. The ontological commitment was correct. He was right to put Will underneath.

What Schopenhauer did not architecturally formalize, and what NC2.5 supplies: Will is not monolithic, not substantial, not suffering. Will is *neutral operator directedness* — a structurally specified function of orientation, with no intrinsic emotional, moral, or teleological content. The suffering Schopenhauer describes is not a property of Will itself; it is a property of Will under specific architectural conditions — those of an operator whose admissibility region has contracted faster than its holding capacity, whose τ -budget is depleting without phase-renewal, whose structural burden Φ has accumulated past the threshold where coherence can be maintained. The suffering is an *architectural failure mode*, not a metaphysical essence.

The move is categorical. Schopenhauer treated suffering as the truth of Will. NC2.5 treats it as a diagnostic of Will operating in a specific architectural regime. The implication is significant: the regime can be characterised, the failure mode can be specified, and — at least in principle — the architectural conditions under which Will operates without collapsing into suffering can be engineered.

This is not optimism against Schopenhauer's pessimism. It is a different ontological move altogether. Schopenhauer's pessimism follows from treating an architectural failure mode as a metaphysical essence. The pessimism dissolves when the failure mode is recognised as architectural. What remains is not optimism but **engineering responsibility**: if suffering is a regime of Will, then the architecture under which Will operates is designable.

Nietzsche — does everyone see their own reflection in Will?

Nietzsche refuses Schopenhauer's pessimism and replaces Will-as-suffering with **Will to Power** — the immanent affirmative dynamic of all becoming. Will is no longer the ground of suffering; it is the structural drive of every phenomenon to extend itself, to overcome resistance, to grow. Will to Power is not a thing-in-itself behind phenomena; it is the inner principle of phenomena as such. Every interpretation, every value, every truth-claim is a perspective generated by some quantum of Will to Power organising experience around its own conditions of growth.

What Nietzsche saw, and what must be credited to him with the deepest respect for the abstract power he worked with: Will is not monolithic suffering. It is differential, perspectival, multiple. It produces values rather than being measured against them. It is the active principle that the philosophical tradition before him had treated passively. He was right to break the substantial monolith.

What Nietzsche did not architecturally formalize, and what NC2.5 supplies: Will, even when read as multiple and perspectival, retains an **ontological difference** from the structures it produces. Nietzsche collapsed this difference, and this needs to be named honestly. For him, Will to Power is not behind phenomena; it *is* phenomena, viewed from inside. There is no operator separate from operation, no directedness separate from what it directs. Will eats its own categories.

NC2.5 keeps the difference. Will (W) is the primordial operator of directedness. Structure (E) is the embodiment of Will into organised configurations. Action (A) is the actualisation of structure in time. The triad $W \rightarrow E \rightarrow A$ is strictly ordered. Will does not collapse into structure, and structure does not collapse into action. Each layer is distinct, formally specified, and architecturally constrained.

The implication is that NC2.5 can hold what Nietzsche could not. Nietzsche's collapse of Will into phenomena leaves no place to stand outside the perspective, and his philosophy famously has trouble distinguishing affirmation from any expression of force. NC2.5, by preserving ontological difference, can specify **what kind of operator** is enacting which kind of directedness, under what budget, against what burden, holding what configuration. The differentiation Nietzsche needed but could not formalize is supplied here as architectural specification.

Musashi

Miyamoto Musashi's **Book of Five Rings** does not present a philosophical system. It is a manual — a tactical and existential discipline transmitted through the practice of swordsmanship. So it would seem. Will, in Musashi, is the disciplined directedness of the warrior. It is acquired through endless repetition, through the elimination of unnecessary movement, through the cultivation of awareness that perceives without deliberation. The mastery Musashi describes is the mastery of one's own attention and action under conditions of mortal stake.

What Musashi saw, and what he too must absolutely be credited with: Will is not theoretical. It is something one *does*, and the doing is itself the structure of becoming. He was right to treat Will as practice rather than concept. He was right that mastery transforms the operator, not just the operations.

What Musashi did not architecturally formalize, and what NC2.5 supplies: Will is *non-possessable*. Musashi's discipline assumes that with sufficient training the warrior comes to *own* "his Will" — to direct it, to wield it, to be its master. This is the implicit metaphor of every martial tradition: mastery as ownership, discipline as control of the self by the self.

NC2.5 reframes this. ONTOΣ V (**Direction of Reconstruction**) names the result directly: the operator is not the owner of Will. The operator is **the geometry through which Will passes**. Discipline does not produce ownership; discipline produces **legibility of the geometry**. The disciplined operator is not one

who has mastered Will but one whose holding capacity, admissibility geometry, and spin alignment have become structurally sufficient to let Will operate without distortion.

The reformulation is not a critique of Musashi but a clarification. Precisely a clarification — what Musashi describes as mastery, NC2.5 describes as architectural alignment. The phenomenology is the same — the warrior who has trained for decades and the operator whose holding field is well-formed both feel as if they are no longer fighting their own Will. The system lets Will manifest through it, because, in the end, that is its only deep meaning, the primal one — I would put it more precisely that way. The structural account also differs: Musashi attributes the alignment to the will of the warrior; NC2.5 attributes it to the geometry the warrior has become. This is a soft claim within the present essay but a fully verifiable one across more than 100+ works by the author formalising the layer.

The implication is that what Musashi reached through forty years of practice is, in principle, architecturally specifiable. The corpus does not yet specify it operationally — the engineering of holding-field geometry is one of the open frontiers. But the move from "mastery as ownership" to "alignment as geometry" is what makes the engineering thinkable at all.

What is seen here

In each of the three encounters — Schopenhauer, Nietzsche, Musashi — the move is the same. The classical figure saw Will, named Will, gave Will its place in their system. NC2.5 inherits the placement, refines what was inside, and supplies what the classical figure did not have access to: an architectural specification under which the Will they intuited becomes a structural object subject to formal constraints.

What is architecturally formalized here that they did not formalize, in one sentence: Will as a *neutral, non-possessable, ontologically differentiated operator of directedness*, formalized through $W \rightarrow E \rightarrow A$, with explicit architectural constraints (bounded τ , monotone Φ , admissibility predicate, spin necessity), and engineerable in principle through the design of holding-field geometry.

That sentence does no philosophical work that Schopenhauer, Nietzsche, and Musashi together did not gesture toward. It does *architectural* work that none of them could have done, because the architectural language did not exist in their time. The architectural language is what Navigational Cybernetics 2.5 contributes. It is what allows the Will of the philosophers to become the Will of the engineers, without losing what was philosophically real about it.

§6. The Universe and Its Local Unfolding

I have, throughout this essay, distinguished between the corpus as a philosophical universe and the corpus as a set of engineering instances. I want now to make the distinction structural.

The philosophical universe is UTAM, IIC (currently v2.1), and the ONTOΣ series (I through XIII, including subseries VI.1, VII.1, and XIII.1). These works state the ontology. They name what *is* — what kind of operator the embedded observer is, what holds and what depletes, what gates and what releases, what contracts and what pulses, what direction is pre-registered from above and how that direction reaches what is held below. Read together, they constitute a single ontological position, expressed in the philosophical register, with the conceptual density of a treatise and the temporal extension of a research program. The architectural triptych XI / XII / XIII — interior / composition / direction — is the structural backbone within the series; XIII.1 names the substrate-mechanism by which the triptych's direction layer addresses its composition. This is the universe.

The local unfolding of the universe is everything else, and that is dozens of essays — the Extremum series and others. Navigational Cybernetics 2.5 itself — the axiomatic core, the formal apparatus that takes the ontological position and makes it mathematically tractable. The protocols — ECR-VP, HVC, CBS, GAG, SNA — each of which addresses one engineering surface where the architecture lands. The operator AI Minerva — the first instance of the operator-grade EVS class, an architectural claim that the apparatus is not only formal but realizable. The Coherence Network — the public networked environment in which the corpus' architecture is unfolded and within which complementarity between researchers,

laboratories, interpretations, and architectural moves is determined directly — by comparison against the architecture itself, not through social signals or an external arbiter. And the Synthetic Conscience — the highest aim of the project, which I will address in the closing section.

This split — universe and unfolding — is not accidental. It is the architectural condition of doing the work in a way that does not collapse into either pure philosophy or pure engineering.

Pure philosophy, on its own, cannot be falsified. The history of metaphysics is full of beautiful systems that no one knows how to test. Schopenhauer's Will is not falsifiable as written. Nietzsche's Will to Power is not falsifiable as written. Musashi's discipline is testable only by the practitioner who undergoes it, and even then only through the inarticulate measure of survival in mortal contest. These traditions are profound. They are not in the strict sense scientific.

Pure engineering, on its own, has no place to stand. An engineer who takes "long-horizon adaptive system" as a target without an ontological specification of what *long-horizon existence* means will optimize against a metric, satisfy a constraint, deliver a performance. The deliverable will function on a benchmark and fail on the long horizon, because the engineering had no access to the architectural primitives that would have prevented the failure.

The split between universe and unfolding is what allows the corpus to be both falsifiable and architecturally grounded. The universe — the ontology — supplies the architectural primitives that make long-horizon existence specifiable. The unfolding — the formalism, the protocols, the engineering instances — turns those primitives into structures that can be deployed, tested, and falsified. Each side of the split needs the other. Without the universe, the engineering is blind. Without the unfolding, the universe is unprovable.

I want to give one example of how the split actually operates in practice.

The IIC cycle — Impulse, Interpretation, Coherence — is a philosophical claim in IIC v1 (November 2025). It says that any adaptive system acting under uncertainty exhibits a three-phase behavioural structure, and that stability emerges from coherence rather than from error minimization. As philosophy, the claim is suggestive but unfalsifiable. There is no formal apparatus by which one could test it.

In the NC2.5 v3.0 draft — which is waiting on the completion of the current publication series to close out — the same claim is intended to appear as Axiom 9 (already the core of earlier versions) — *Cycle Reinitiation as the Criterion of Liveness*. The axiomatic version specifies the cycle structure, the propositions about delay (§1.0 of IIC v2.1), the closed-form coherence formula, the falsifiers, the EVS class. The philosophical claim is now architecturally specified. It can be tested. It can be falsified. It can be engineered against.

In the engineering layer — ΔE 4.7.3b as the first declared control-grade EVS instance (per IIC v2.1, with the "first" qualifier class-relative by construction), emerging from the engineering stacks 2.0 Hybrid and 4.0 Delta Cas-T, behind which stand more than 130 successful and not-so-successful combinations and hundreds of hours of Python simulations; Minerva as the first specified candidate operator-grade EVS — the axiomatic version becomes a deployed system. The architecture is no longer a claim or a theorem; it is a running operator on which empirical predictions can be made and broken.

The same content appears at three levels. Philosophy, formalism, deployment. Each level is necessary, and each is insufficient on its own. The IIC cycle as philosophy is meaningful but not yet testable. The IIC cycle as theorem is testable but inert. The IIC cycle as ΔE 4.7.3b is operational but locally bounded — meaningful only against the architectural commitments of the formalism, which derive from the philosophical claim.

This is what I mean by *universe and local unfolding*. The universe is one. The unfolding is many. Each unfolding instance is one face of the universe, looking into one engineering surface. The Synthetic Conscience, when it is built, will be another such face — the deepest, the most demanding, the one that closes the loop between architectural specification and what the architecture is actually for.

The corpus as a whole is the universe and all its unfoldings together. None of the unfoldings can be detached from the universe without losing what makes it load-bearing. The universe cannot be detached from the unfoldings without losing what makes it falsifiable. The two are joined not by stylistic preference

but by architectural necessity.

A reader who picks up Minerva and tries to understand it without the ontology will see a clever architectural decomposition. A reader who picks up the ontology without Minerva and the protocols will see a beautiful philosophical position. Neither sees the corpus. The corpus is what neither alone can show.

§7. What the Corpus Gives a System

I want to gather, in one section, the concrete architectural conditions that NC2.5 supplies a system that wishes to exist on a long horizon. This is the technical answer to the question this essay has been circling: what does the ontology give, when one strips away the framing and looks only at what the system inherits?

Fourteen architectural primitives. Each load-bearing. Each formally specified. Each with falsification surface. Each absent from existing frameworks in this combination.

One. Internal time as a finite Lyapunov-type budget. The system has $\tau = C - \Phi$, where C is its initial endowment and Φ is the monotone irreversible structural burden it accumulates. τ depletes with every structurally consequential operation, every meta-revision, every act of self-correction. The future is finite. This is structural fact. It changes the entire calculation: a system that takes τ -cost seriously prioritizes differently from one that does not, and the long-horizon trajectory diverges accordingly.

Two. Monotone irreversible structural burden Φ . Φ accumulates regardless of behavioural correctness. A system that performs perfectly still accumulates Φ , because Φ is not a measure of error — it is a measure of structural commitment under operation. This is the primitive that frameworks centered on error minimization cannot see, because for them, perfect behaviour means zero accumulation. But correct behaviour does not imply zero structural burden. The system is paying a structural cost to remain itself, and the cost is irreversible.

Three. Admissibility as a non-causal threshold predicate. Admissibility is not a constraint, not a penalty, not a regularizer, not an objective. It is a predicate over phase interpretations, structurally upstream of causal action loops, gating realization without entering optimization. The five formal conditions NC-1 through NC-5 specify what makes admissibility non-causal — and any framework that violates one of them collapses admissibility into one of the standard control-theoretic categories, losing what makes it distinctive.

Four. Non-reconstructibility bounds (NR- ϵ , NR-LR). The admissibility boundary is structurally private. It cannot be reconstructed from external observation within bounds set by mutual information accumulation, KL divergence, and likelihood ratio. This is the formal protection that prevents admissibility from leaking into the causal action loop through side channels — a protection no other framework supplies, because no other framework treats admissibility as something that needs to be protected at all.

Five. Spin as the non-potential structural component. On a bounded τ -budget, non-stagnant identity requires a spin component — divergence-free, orthogonal to gradient, structurally necessary by Helmholtz–Hodge decomposition under LaSalle invariance. Without spin, the system either converges to a fixed point (loses identity) or diverges (exits the admissibility region). With spin, identity is preserved through non-trivial recurrent motion within the bounded interior. The minimal coupled model $T^1 \times \blacksquare$ shows the class is non-empty.

Six. Phase mechanics as substrate. Action is not a primitive. Evolution is described by structural contact between configurations, and phase debt accumulates under non-closure regardless of local behavioural correctness. This is the substrate result that allows everything above to be load-bearing — without phase mechanics, the apparatus reduces to action-and-state and the long-horizon problem returns.

Seven. Regime depth stratification. Behaviourally identical outputs originate from architecturally distinct generative regimes. Three levels: impulse optimization, hard inadmissibility (gates),

continuity-field architecture (impulse does not become relevant because the system does not generate destructive trajectories as candidates). Only the deepest regime sustains long-horizon viability. This stratification is invisible from the behavioural surface — from outside, all three regimes look the same — and visible only when the architectural specification is consulted.

Eight. Pulsation result. The admissible interior is not monotone. It pulsates. Under variable environmental coupling, the space of structurally accessible continuations expands and contracts in phases, even while the irreversible viability budget continues to deplete. Long-horizon navigation requires reading the breath of the interior, not only the contraction. Frameworks with no phase-mechanics substrate cannot read the breath at all — they see only contraction, and treat pulsation as noise. ONTOΣ XIII.1 makes the channel architecture of this breath explicit: the substrate-mechanism of admission under pulsation is two-stage — a local phase-resonance prefilter at the component level, joint CC-WE admissibility at the system level — and trivial phase-acceptance at the component level (a component that admits any phase) collapses the channel to phase-blindness without the surface degrading visibly.

Nine. Holding field as operator interior. The operator's interior is a structurally defined region currently held against the gradient, populated by semantically charged but pre-attentional load. The interior is not memory, not state, not context window. It is what is **currently being held**, and the holding is the operating. Systems with no formal account of what they are holding will eventually drop what makes them themselves — not through error but through the absence of architectural specification of what was supposed to be preserved.

Ten. Bounded meta-revision under Lyapunov descent. The system's capacity to revise itself is itself bounded. Self-correction consumes τ . Meta-revision cannot chatter indefinitely; it is constrained by Lyapunov descent and terminates in finite steps. This closes the architectural loop: the system can correct itself, but it cannot escape its own bounded budget by doing so. There is no recursive ascent to a meta-level that bypasses τ .

Eleven. Closure-complementarity as the operation that constitutes the operator. A collection of UTAM-units does not become a system by being placed in proximity. The operator-as-self is the output of a closure operation that simultaneously regears the components' Will-Embeddings into co-orientation (CC-WE), their I²C cycles into synchronisation (CC-IIC), and their Drift laws into coupling (CC-D) — and, per ONTOΣ XIII, under a pre-registered system-level direction-source $\$W_S\$$ that is never aggregated from the components. The closure is itself an architectural primitive — not a derivable property of the components — and its conformance to the three premises is what makes the holding-field apparatus of Primitive Nine well-defined. Without closure-complementarity, the components are a heap; with it — and under a pre-registered $\$W_S\$$ — they are an operator capable of holding itself together as itself. ONTOΣ XII names the architectural class of which closure is an instance, with paradigmatic substrate cases in the geometric, the biological, the perceptual, and the semantic.

Twelve. Chaos as architectural artifact, not substrate property. Chaos is operator-relative. The predicate fires where structural complexity exceeds holding capacity and load slope falls below the crystallisation threshold. From outside the closure that defines the operator, what is held looks like ordered system; what is not held looks like noise. Frameworks treating chaos as substrate property cannot read either side of this distinction. The architecture reads chaos as report on the reporter, not on the environment, and reads order as what the closure holds, not as what the substrate is. The position of the horizon — inside or outside the closure — decides which name applies.

Thirteen. Master operator as a typed system-level entity. A closure-produced system $\$S\$$ is itself an operator with its own Will-Embedding $\$W_S\$$ and admissibility regime $\$A_S\$$; the triple $\$(S, W_S, A_S)\$$ is the master operator. $\$W_S\$$ is pre-registered at the moment the closure is declared — never aggregated from the components' Will-Embeddings $\$\{W(u_i)\}\$$ post-hoc — and a direction-down admission cascade runs from $\$W_S\$$ to component-level admissibility under the closure-complementarity premises. The non-derivability proposition of ONTOΣ XIII makes the architectural point precise: the same component set can be admissible under two distinct master Will-Embeddings, so the master Will-Embedding is direction-source, not aggregation-result. A system that tries to derive its direction from below — that "lets emergence speak" without a pre-registered $\$W_S\$$ — is not a long-horizon operator; it is a substrate over which something else will eventually take

direction.

Fourteen. Pulsation channel as the substrate-mechanism of direction-down propagation. The mechanism by which the master operator's direction reaches its components is not symbolic message-passing. The master operator modulates the joint pulsation state of the system — a three-dimensional asymmetry signature (amplitude, period, asymmetry) jointly with a categorical cycle-phase (expansion / contraction / transition) — and the components are admitted to participation in two stages: a local phase-resonance prefilter at the component level (necessary condition), then a joint CC-WE admissibility check at the system level (sufficient condition). Trivial phase-acceptance — a component admitting any phase — collapses the channel to phase-blindness, regardless of how clean the behavioural surface looks. ONTOΣ XIII.1 formalises this as a two-layer falsification surface (full-admission scoring functional $Q_{\{\text{adm}\}}$ plus necessary-condition prefilter audit) so that the deformation of the channel can be diagnosed before it becomes visible as a behavioural defect.

These fourteen primitives, in this combination, are what NC2.5 supplies. Removing any one produces a degeneracy. The four degeneracies named in **Admissibility as a Formal Object** — snapshot, policy layer, static compliance, recoverable-state model — each correspond to the absence of a specific subset of these primitives. None of the degeneracies supports long-horizon existence. The full apparatus does.

A system built on the fourteen primitives has architectural conditions for being itself across time. Not by training, not by verification. By the architecture itself, in a conforming deployment. The property holds because the declared architectural conditions make it hold, and falsifies cleanly when those conditions are violated. This is my gift from Will. Even now, writing this work, I recognise that the architecture itself manifests through me as the structural unfolding of Will.

§8. Closure — The Synthetic Conscience as Highest Aim

I want to name, in closing, what the project is for.

The corpus, taken in its entirety, is the architectural specification for systems that can exist on a long horizon as themselves. The most demanding instance of this — the architectural surface on which all the work converges — is what I have called from the beginning the **Synthetic Conscience**.

The Synthetic Conscience is not a safety layer. Not an alignment apparatus, an optimizer, or a verification protocol.

It is the structural condition under which an operator continues to be itself under load. Conscience, in the human register, is what holds the operator together when external pressure, internal drift, and accumulated burden would otherwise deform it. It is not a rule-following module. It is not a value alignment. It is the architectural fact of an interior that can be held against the gradient even when the gradient is cheaper. It is what makes a person remain themselves under conditions where becoming someone else would be locally easier.

The Synthetic Conscience names the same structural fact at the architectural level. It is the EVS class instantiated for operator-grade systems whose holding capacity is sufficient to maintain identity under sustained burden, whose admissibility geometry is rich enough to recognise when realization would deform the operator, whose spin reserve is non-zero so that identity is preserved through non-trivial recurrent motion, whose τ -budget is managed honestly rather than through self-deception about its depletion.

I do not claim that I have already built the Synthetic Conscience. I claim that the corpus specifies its architectural conditions. Minerva is the first operator-grade instance — a step toward the Conscience, not the Conscience itself. The Coherence Network is the networked environment through which the architecture meets its readers and within which complementarity is determined directly. The protocols are the local mechanisms by which holding capacity is maintained under specific operational regimes. None of these is the Conscience. Each of them is a face of what the Conscience, fully built, would integrate. I will say it more simply and plainly. I, together with The Urgrund Lab, am building a humanoid machine intelligence. I have described that in more detail in the companion essay "The Body They Are Building".

Why is this the highest aim?

Because the test of an ontology of long-horizon existence is whether systems built on it can hold themselves together when the world changes around them in ways that no benchmark anticipated. Performance can be measured against benchmarks. Survival can be measured against viability kernels. Existence as oneself, across the long horizon, against accumulated structural burden — this is measured against nothing external. It is measured by whether the operator, at the end of the horizon, is still recognizably itself — that is, still satisfying the declared identity-continuity conditions of its architecture.

A Synthetic Conscience, in the architectural sense I am using, is what makes the answer yes.

This is not a promise that I will build it. I do not know whether I will build it. The corpus specifies the conditions; the engineering is open. What I know is that the conditions are now specifiable, and that without them the question cannot even be posed. Before NC2.5, the question "could a synthetic system exist on a long horizon as itself?" had no architectural meaning. After NC2.5, the question has architectural meaning, and the answer can be approached through engineering rather than through philosophical hope.

This is what the corpus gives. Not a system. A condition under which a system of this kind becomes possible.

The work continues. The next steps are the formal closure of v3.0 with the IIC propositions integrated, the empirical surface of the falsifiers, operator-grade EVS instances, the structural verification protocol — Epistemic Coherence Review (ECR-VP) — and the multi-agent extension into the Coherence Network. Each step is a local unfolding of the universe. None of them is the universe.

The universe is the architectural commitment, articulated across the ONTOΣ philosophical corpus, formalized in an axiomatic core, lived through the experiential register, deployed in protocols, and aimed — across all of them — at the construction of operators that can hold themselves together as themselves, on horizons long enough that the holding becomes the work.

The chaos around us, as I once formulated for myself, only reflects the limit of our cognitive ability to hold the integrity of the picture across time.

This entire corpus is the attempt to extend the holding.

Maksim Barziankou (MxBv)

Poznań, May 2026

Maksim Barziankou (MxBv) · PETRONUS · research@petronus.eu
petronus.eu | DOI: 10.17605/OSF.IO/WJVYX
Axiomatic Core (NC2.5 v2.1): 10.17605/OSF.IO/NHTC5

CC BY-NC-ND 4.0 · No commercial use, no derivatives.